

Arithmétique flottante et analyse d'erreur (AFAE)

Lecture 1: summation

Theo Mary (CNRS)

theo.mary@lip6.fr

<https://perso.lip6.fr/Theo.Mary/>

M2 course at Sorbonne Université,
2025–2026

$$\sum_{i=1}^n x_i$$

Introduction

Dealing with accumulation

Dealing with cancellation

Adaptive precision summation

Conclusion

Why summation is important

$$y = \sum_{i=1}^n x_i \quad \dots \text{an ubiquitous and fundamental task!}$$

- **Dot products:**

$$a, b \in \mathbb{R}^n \Rightarrow a^T b = \sum_{i=1}^n a_i b_i$$

- **Matrix–vector products:**

$$A \in \mathbb{R}^{m \times n}, \ b \in \mathbb{R}^n \Rightarrow (Ab)_j = \sum_{i=1}^n a_{ji} b_i, \quad j = 1: m$$

- **Matrix–matrix products:**

$$A \in \mathbb{R}^{m \times n}, \ B \in \mathbb{R}^{n \times p} \Rightarrow (AB)_{jk} = \sum_{i=1}^n a_{ji} b_{ik}, \quad j = 1: m, k = 1: p$$

- **Gaussian elimination (LU factorization):**

$$A \in \mathbb{R}^{n \times n}, \ A = LU \Rightarrow \begin{cases} \ell_{jk} &= \left(a_{jk} - \sum_{i=1}^{k-1} \ell_{ji} u_{ik} \right) / u_{kk} \\ u_{kj} &= a_{kj} - \sum_{i=1}^{k-1} \ell_{ki} u_{ij} \end{cases}$$

Accumulation of rounding errors

Summation suffers from the accumulation of rounding errors

Standard model of FP arithmetic:

$$\text{fl}(x \text{ op } y) = (x \text{ op } y)(1 + \delta), \quad |\delta| \leq u, \text{ for op} \in \{+, -, \times, \div\}$$

Consider the computation of $y = \sum_{i=1}^n x_i$ by recursive summation:

$$\begin{aligned} y_2 &= x_1 + x_2 \quad \Rightarrow \quad \hat{y}_2 = (x_1 + x_2)(1 + \delta_1) \\ y_3 &= \hat{y}_2 + x_3 \quad \Rightarrow \quad \hat{y}_3 = (\hat{y}_2 + x_3)(1 + \delta_2) \\ &\quad = (x_1 + x_2) \underbrace{(1 + \delta_1)(1 + \delta_2)}_{\delta_1 \text{ and } \delta_2 \text{ accumulate!}} + x_3(1 + \delta_2) \\ y_4 &= \dots \text{etc.} \end{aligned}$$

How can we measure the accumulated effect of all rounding errors?

Forward and backward errors

- Let $y = f(x)$ be computed in finite precision and let \hat{y} be the computed result
- **Forward error** analysis measures

$$|\hat{y} - y| \text{ (absolute)} \quad \text{or} \quad \frac{|\hat{y} - y|}{|y|} \text{ (relative)}$$

- **Backward error** analysis computes the smallest perturbation Δx such that

$$\hat{y} = f(x + \Delta x)$$

and measures $|\Delta x|$ (absolute) or $|\Delta x|/|x|$ (relative).

- Backward error analysis recasts the rounding errors as perturbations of the input data
- An algorithm is **backward stable** if it yields a small backward error, where “small” usually means $O(u)$

Forward and backward errors for summation

- **Forward error**

$$\eta_{\text{fwd}} = \frac{|\hat{y} - y|}{|y|}$$

- **Backward error**

$$\eta_{\text{bwd}} = \min \left\{ \varepsilon > 0 : \exists \delta x_i, \hat{y} = \sum_{i=1}^n x_i + \delta x_i, |\delta x_i| \leq \varepsilon |x_i| \right\}.$$

Two questions:

- Find a **formula** for η_{bwd}
- Find **bounds** for η_{bwd} and η_{fwd} when \hat{y} is computed in floating-point arithmetic

Formula for backward error

$$\eta_{\text{bwd}} = \min \left\{ \varepsilon > 0 : \exists \delta x_i, \hat{y} = \sum_{i=1}^n x_i + \delta x_i, |\delta x_i| \leq \varepsilon |x_i| \right\}.$$

We have the formula

$$\eta_{\text{bwd}} = \frac{|\hat{y} - y|}{\sum_{i=1}^n |x_i|}.$$

Proof:

- $\frac{|\hat{y} - y|}{\sum_{i=1}^n |x_i|} \leq \eta_{\text{bwd}}$
- $\eta_{\text{bwd}} \leq \frac{|\hat{y} - y|}{\sum_{i=1}^n |x_i|}$ (using $\delta x_i = (\hat{y} - y) \frac{|x_i|}{\sum_{i=1}^n |x_i|}$)

Formula for conditioning

As a result we also obtain the formula

$$\kappa = \frac{\eta_{\text{fwd}}}{\eta_{\text{bwd}}} = \frac{\sum_{i=1}^n |x_i|}{\left| \sum_{i=1}^n x_i \right|}.$$

- κ is large if $\sum |x_i| \gg |\sum x_i| \Rightarrow \text{cancellation}$

Backward error analysis

$$\begin{aligned}y_2 &= x_1 + x_2 \\ \Rightarrow \hat{y}_2 &= (x_1 + x_2)(1 + \delta_1) = x_1(1 + \delta_1) + x_2(1 + \delta_1) \\ y_3 &= \hat{y}_2 + x_3 \\ \Rightarrow \hat{y}_3 &= (\hat{y}_2 + x_3)(1 + \delta_2) \\ &= x_1(1 + \delta_1)(1 + \delta_2) + x_2(1 + \delta_1)(1 + \delta_2) + x_3(1 + \delta_2) \\ \dots \\ \Rightarrow \hat{y}_n &= \sum_{i=1}^n \left[x_i \prod_{k=k_i}^n (1 + \delta_k) \right]\end{aligned}$$

Backward error analysis

$$\begin{aligned}y_2 &= x_1 + x_2 \\ \Rightarrow \hat{y}_2 &= (x_1 + x_2)(1 + \delta_1) = x_1(1 + \delta_1) + x_2(1 + \delta_1) \\ y_3 &= \hat{y}_2 + x_3 \\ \Rightarrow \hat{y}_3 &= (\hat{y}_2 + x_3)(1 + \delta_2) \\ &= x_1(1 + \delta_1)(1 + \delta_2) + x_2(1 + \delta_1)(1 + \delta_2) + x_3(1 + \delta_2) \\ \dots \\ \Rightarrow \hat{y}_n &= \sum_{i=1}^n \left[x_i \prod_{k=k_i}^n (1 + \delta_k) \right]\end{aligned}$$

Worst-case fundamental lemma

Let δ_k , $k = 1 : n$, such that $|\delta_k| \leq u$ and $nu < 1$. Then

$$\prod_{k=1}^n (1 + \delta_k) = 1 + \theta_n, \quad |\theta_n| \leq \gamma_n := \frac{nu}{1 - nu}.$$

General worst-case bound

General algorithm

$$\mathbb{S} = \{x_1, \dots, x_n\}$$

Repeat

 Choose any pair $(x_i, x_j) \in \mathbb{S}^2$ ($i \neq j$)

$$\mathbb{S} \leftarrow \mathbb{S} \setminus \{x_i, x_j\}$$

$$\mathbb{S} \leftarrow \mathbb{S} \cup \{x_i + x_j\}$$

until $\mathbb{S} = \{y\}$

No matter the summation order we have the bound

$$\eta_{\text{bwd}} \leq \gamma_{n-1} = (n-1)u + O(u^2)$$

Summary

Consider the computation

$$y = \sum_{i=1}^n x_i$$

In floating-point arithmetic, the forward error η_{fwd} is bounded by

$$\eta_{\text{fwd}} \leq \eta_{\text{bwd}} \kappa, \quad \eta_{\text{bwd}} \leq \gamma_{n-1} = (n-1)u + O(u^2), \quad \kappa = \frac{\sum |x_i|}{|\sum x_i|}$$

Thus η_{fwd} can be large when

- The **unit roundoff u** is large (**low precision**)
- The **dimension n** is large (**accumulation**)
- The **condition number κ** is large (**cancellation**)

Introduction

Dealing with accumulation

Dealing with cancellation

Adaptive precision summation

Conclusion

No matter the summation order we have the bound

$$\eta_{\text{bwd}} \leq \gamma_{n-1} = (n-1)u + O(u^2)$$

⇒ However, for specific orders, we can get much better bounds, and much smaller errors!

Summation tree

Given a summation order to compute $y = \sum_{i=1}^n x_i$, we define its associated summation tree as a **binary tree** such that:

- the n **leaf nodes** are the n summands x_i ;
- any **inner node** is equal to the sum of its two children
- the **root node** is the final sum y

Example: recursive summation is a **comb tree**

Tree-based summation

- For any summation tree, we have the bound:

$$\eta_{\text{bwd}} \leq \gamma_h = hu + O(u^2)$$

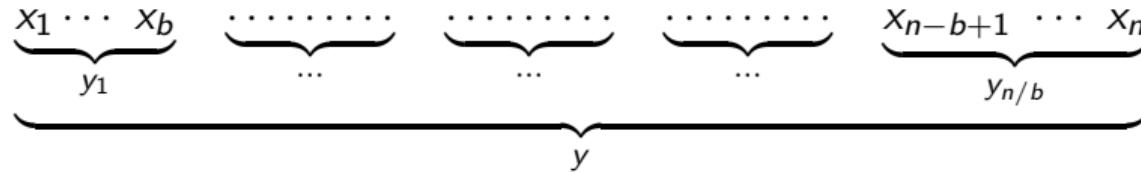
where h is the height of the tree

- The minimal bound is therefore attained for a **balanced binary tree**, for which $h = \lceil \log_2 n \rceil$. This is called **pairwise summation**.
- While it achieves the minimal bound, pairwise summation is not efficient on modern computers.

Blocked summation

Blocked summation algorithm:

```
for  $i = 1 : n/b$  do
    Compute  $y_i = \sum_{j=(i-1)b+1}^{ib} x_j.$ 
end for
Compute  $y = \sum_{i=1}^{n/b} y_i.$ 
```



- Widely used in NLA libraries (BLAS, LAPACK, ...)
- $\eta_{\text{bwd}} \leq \gamma_h$ with $h = b + n/b - 2$
- With optimal $b = \sqrt{n}$: $h = 2(\sqrt{n} - 1)$

Blocked summation

for $i = 1 : n/b$ **do**

 Compute $y_i = \sum_{j=(i-1)b+1}^{ib} x_j$.

end for

Compute $y = \sum_{i=1}^{n/b} y_i$.

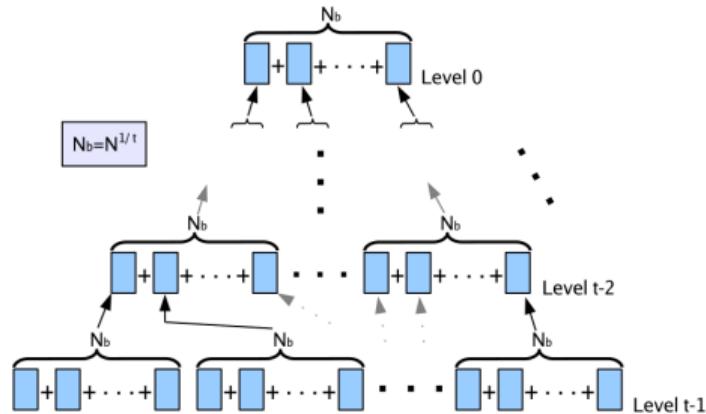
$$\hat{y}_i = \sum_{j=(i-1)b+1}^{ib} \left[x_j \underbrace{\prod_{k=k_j}^b (1 + \delta_k^{(i)})}_{\text{at most } b - 1 \text{ terms}} \right]$$

$$\hat{y} = \sum_{i=1}^{n/b} \left[\hat{y}_i \underbrace{\prod_{k=k'_i}^{n/b} (1 + \delta'_k)}_{\text{at most } n/b - 1 \text{ terms}} \right]$$

$$= \sum_{j=1}^n \left[x_j \underbrace{\prod_{k=k_j}^b (1 + \delta_k^{(i)})}_{\text{at most } b + n/b - 2 \text{ terms}} \underbrace{\prod_{k=k''_j}^{n/b} (1 + \delta'_k)}_{\text{at most } n/b - 1 \text{ terms}} \right]$$

Superblock summation

- **Superblocked summation:** tree summation with t levels, block size at level t :
 $b_t = n^{1/t}$
 - $t = 1 \Rightarrow$ standard recursive summation
 - $t = 2 \Rightarrow$ optimal blocked summation
 - $t = \log_2 n \Rightarrow$ pairwise summation
 - $\eta_{\text{bwd}} \leq \gamma_h$ with $h = t(n^{1/t} - 1)$
 - Castaldo et al. (2009)



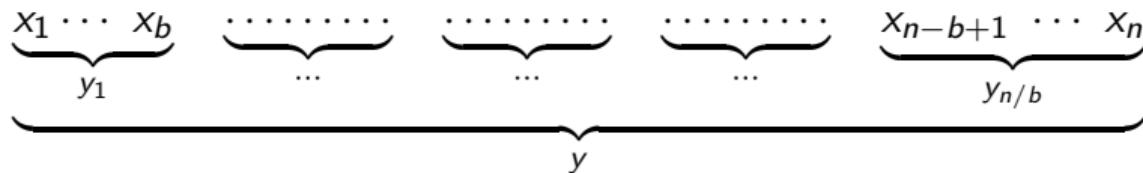
Fast Accurate Blocked summation algorithm (FABsum) [Blanchard, Higham, M. \(2020\)](#)

for $i = 1 : n/b$ **do**

 Compute $y_i = \sum_{j=(i-1)b+1}^{ib} x_j$ with **FastSum**.

end for

Compute $y = \sum_{i=1}^{n/b} y_i$ with **AccurateSum**.

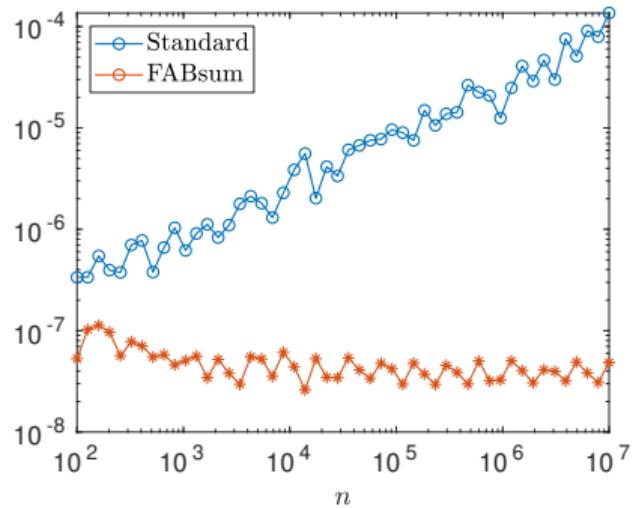


- Cost: $C(n, b) = \frac{n}{b}C_f(b) + C_a(\frac{n}{b}) \approx C_f(n) + \frac{1}{b}C_a(n)$
- Error: $\epsilon(n, b) = \epsilon_f(b) + \epsilon_a(n/b) + \epsilon_f(b)\epsilon_a(n/b)$
 \Rightarrow If $\epsilon_a(p) = pu^2$ (recursive summation in precision u^2), then $\epsilon(n, b) = bu + O(u^2)$ is independent of n to first order

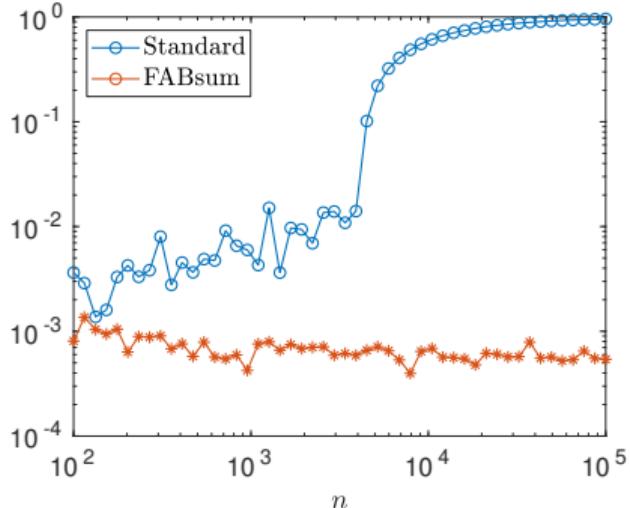
FABsum: numerical results

Backward error for summing random uniform $[0, 1]$ data

fp32



fp16



Blocked summation: implementation remark

```
for i = 1: n/b do
    Compute  $y_i = \sum_{j=(i-1)b+1}^{ib} x_j$ .
end for
Compute  $y = \sum_{i=1}^{n/b} y_i$ .
```

- If implemented as is, requires storing n/b intermediate y_i values, which requires extra memory and is likely to slow down computation
- Better to implement as follows:

```
 $y = 0$ 
for i = 1: n/b do
    Compute  $z = \sum_{j=(i-1)b+1}^{ib} x_j$ .
    Compute  $y = y + z$ 
end for
```

Introduction

Dealing with accumulation

Dealing with cancellation

Adaptive precision summation

Conclusion

Fast2Sum

$$[x, y] = \text{Fast2Sum}(a, b)$$

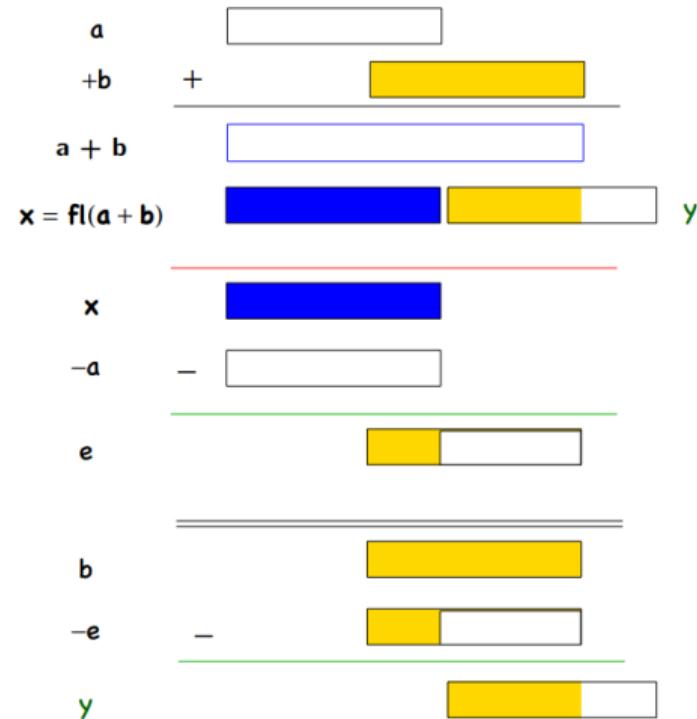
Input: $a, b \in \mathbb{F}$ such that $|a| \geq |b|$

Output: $x = \text{fl}(a + b), y \in \mathbb{F}$ such that
 $x + y = a + b$

$$x = a + b$$

$$e = x - a$$

$$y = b - e$$



Fast2Sum

$$[x, y] = \text{2Sum}(a, b)$$

Input: $a, b \in \mathbb{F}$ such that $|a| \geq |b|$

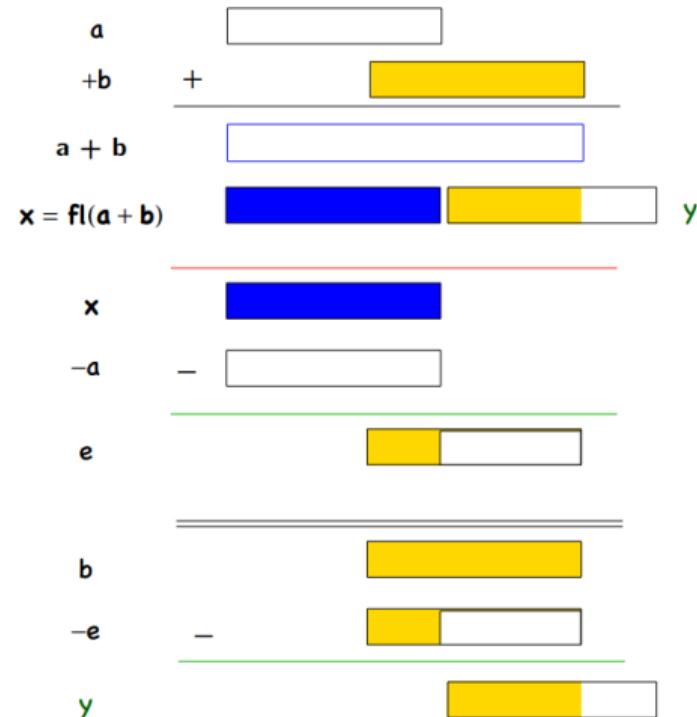
Output: $x = \text{fl}(a + b), y \in \mathbb{F}$ such that
 $x + y = a + b$

$$x = a + b$$

$$e = x - a$$

$$y = (a - (x - e)) + (b - e)$$

Can remove $|a| \geq |b|$ restriction at the cost
of 3 extra flops



Kahan's summation (compensated summation)

Input: $x_i \in \mathbb{F}$, $i = 1: n$

Output: $y \approx \sum_{i=1}^n x_i$

$y = 0, z = 0$

for $i = 1: n$ **do**

$t = x_i + z$

$[y, z] = \text{Fast2Sum}(y, t)$

end for

- Kahan's summation reinjects the errors at each step in the sum
- It satisfies the bound $\eta_{\text{bwd}} \leq 2u + O(nu^2)$ (proof is quite complicated)

Kahan–Babuška's variant

Input: $x_i \in \mathbb{F}$, $i = 1 : n$

Output: $y \approx \sum_{i=1}^n x_i$

$$y = 0, e = 0$$

for $i = 1 : n$ **do**

$$[y, z] = \text{Fast2Sum}(y, x_i)$$

$$e = e + z$$

end for

$$y = y + e$$

- Kahan–Babuška's variant accumulates the errors separately and adds them to the sum at the end
- It satisfies the bound $\eta_{\text{bwd}} \leq 2u + n^2 u^2$ (proof is still quite complicated)

Kahan–Babuška's variant with branching: Neumaier's variant

Input: $x_i \in \mathbb{F}$, $i = 1 : n$

Output: $y \approx \sum_{i=1}^n x_i$

```
 $y = 0, e = 0$ 
for  $i = 1 : n$  do
    if  $|y| \geq |x_i|$  then
         $[y, z] = \text{Fast2Sum}(y, x_i)$ 
    else
         $[y, z] = \text{Fast2Sum}(x_i, y)$ 
    end if
     $e = e + z$ 
end for
 $y = y + e$ 
```

- Remember that $\text{Fast2Sum}(a, b)$ is only exact if $|a| \geq |b|$
- It satisfies the bound $\eta_{\text{fwd}} \leq u + n^2 \kappa u^2$

Kahan–Babuška's variant with branching: Neumaier's variant

Input: $x_i \in \mathbb{F}$, $i = 1 : n$

Output: $y \approx \sum_{i=1}^n x_i$

$y = 0$, $e = 0$

for $i = 1 : n$ **do**

if $|y| \geq |x_i|$ **then**

$[y, z] = \text{Fast2Sum}(y, x_i)$

else

$[y, z] = \text{Fast2Sum}(x_i, y)$

end if

$e = e + z$

end for

$y = y + e$

Proof: thanks to the branching, Fast2Sum is now error-free so before the final $y + e$ addition, we have

$$\sum_{i=1}^n x_i = y + \sum_{i=1}^n z_i$$

and after it we thus obtain

$$\hat{y} = y + \hat{e} + \delta, \quad |\delta| \leq u|y + \hat{e}|$$

$$\hat{e} = \sum_{i=1}^n z_i + \Delta e, \quad |\Delta e| \leq \gamma_{n-1} \sum_{i=1}^n |z_i|$$

$$\leq \gamma_{n-1}^2 \sum_{i=1}^n |x_i|$$

- Remember that $\text{Fast2Sum}(a, b)$ is only exact if $|a| \geq |b|$
- It satisfies the bound $\eta_{\text{fwd}} \leq u + O(n^2)\kappa u^2$

Avoiding branching: Sum2

Input: $x_i \in \mathbb{F}$, $i = 1 : n$

Output: $y \approx \sum_{i=1}^n x_i$

```
 $y = 0, e = 0$ 
for  $i = 1 : n$  do
     $[y, z] = \text{2Sum}(y, x_i)$ 
     $e = e + z$ 
end for
 $y = y + e$ 
```

- Replacing `Fast2Sum` by `2Sum` avoids branching but requires more flops

Sum2: variant with vector overwriting

Input: $x_i \in \mathbb{F}$, $i = 1 : n$

Output: $y \approx \sum_{i=1}^n x_i$

for $i = 2 : n$ **do**

$$[x_i, x_{i-1}] = 2\text{Sum}(x_i, x_{i-1})$$

end for

$$y = \left(\sum_{i=1}^{n-1} x_i \right) + x_n$$

- Let x' be the overwritten vector (for clarity of notation)
- x'_1, \dots, x'_{n-1} are overwritten by the errors and x'_n by the floating-point evaluation of the sum $\sum_{i=1}^n x_i$
- Hence $|x'_i| \leq O(n)u \sum_{i=1}^n |x_i|$ and $|x'_n| \leq |\sum_{i=1}^n x_i| + O(n)u \sum_{i=1}^n |x_i|$
- Therefore the condition number of x' has been reduced by a factor $O(n^2)u$:

$$\frac{\sum_{i=1}^n |x'_i|}{|\sum_{i=1}^n x'_i|} = \frac{\sum_{i=1}^n |x'_i|}{|\sum_{i=1}^n x_i|} \leq \frac{O(n^2)u \sum_{i=1}^n |x_i| + |\sum_{i=1}^n x_i|}{|\sum_{i=1}^n x_i|} = O(n^2)u\kappa + 1$$

SumK: K -fold compensated summation

Input: $x_i \in \mathbb{F}$, $i = 1 : n$

Output: $y \approx \sum_{i=1}^n x_i$

for $k = 1 : K - 1$ **do**

for $i = 2 : n$ **do**

$[x_i, x_{i-1}] = 2\text{Sum}(x_i, x_{i-1})$

end for

end for

$$y = \left(\sum_{i=1}^{n-1} x_i \right) + x_n$$

- After K iterations, the condition number of x' is $O((nu)^K)\kappa$
- Hence we have the bound $\eta_{\text{fwd}} \leq u + O(n^k)\kappa u^k$
 - ❑ Ogita, Rump, Oishi (2005)
- However, we do not know κ , so how do we know when to stop? \Rightarrow **AccSum**
 - ❑ Rump, Ogita, Oishi (2008)

Distillation methods

$$\sum_{i=1}^n x_i \xrightarrow{\text{distillation}} \sum_{i=1}^n d_i, \quad \text{where } \kappa(d_i) \ll \kappa(x_i)$$

- Goal: distill until $\sum_i d_i$ can be accurately evaluated in floating-point arithmetic
- Higher $\kappa \Rightarrow$ more iterations!

Condensation methods

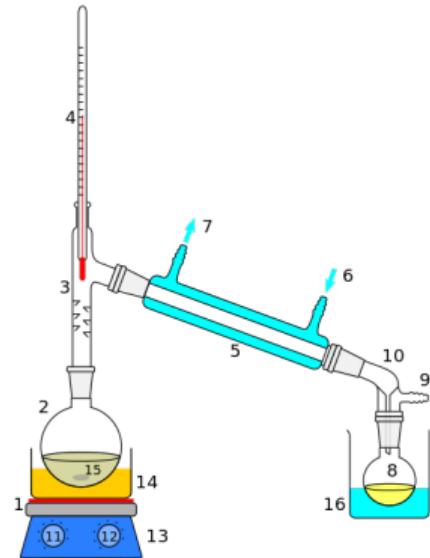
$$\sum_{i=1}^n x_i \xrightarrow{\text{condensation}} \sum_{i=1}^m c_i \xrightarrow{\text{distillation}} \sum_{i=1}^m d_i, \quad \text{where } m \ll n \text{ and } \kappa(d_i) \ll \kappa(x_i)$$

- Goal: reduce the number of summands before applying the costly distillation

Condensation methods

$$\sum_{i=1}^n x_i \xrightarrow{\text{condensation}} \sum_{i=1}^m c_i \xrightarrow{\text{distillation}} \sum_{i=1}^m d_i, \quad \text{where } m \ll n \text{ and } \kappa(d_i) \ll \kappa(x_i)$$

- Goal: reduce the number of summands before applying the costly distillation



Conceptual algorithm

$$\mathbb{S} = \{x_1, \dots, x_n\}$$

Repeat for all pairs $(x_i, x_j) \in \mathbb{S}^2$ ($i \neq j$) such that $x_i + x_j$ is exact

$$\mathbb{S} \leftarrow \mathbb{S} \setminus \{x_i, x_j\}$$

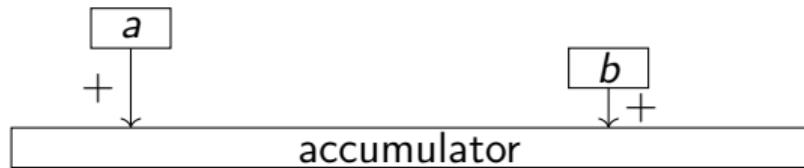
$$\mathbb{S} \leftarrow \mathbb{S} \cup \{x_i + x_j\}$$

until no such pair remains

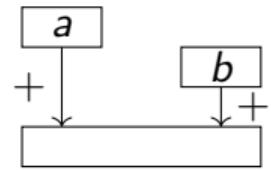
Distill \mathbb{S}

- Can we easily determine when $x_i + x_j$ is exact?
- Can we bound the maximum number of leftover summands?

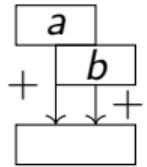
Demmel–Hida method



Demmel–Hida method

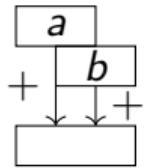


Demmel–Hida method



Consider arithmetic with f -bit mantissa and
 e -bit exponent ($e = 11$ for fp64).

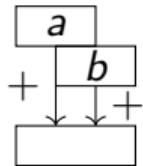
Demmel–Hida method



Consider arithmetic with f -bit mantissa and e -bit exponent ($e = 11$ for fp64).

- One big accumulator: **Kulisch** method
... need one accumulator of $2^e + \log_2 n$ bits

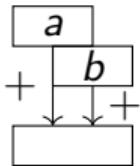
Demmel–Hida method



Consider arithmetic with f -bit mantissa and e -bit exponent ($e = 11$ for fp64).

- One big accumulator: **Kulisch** method
... need one accumulator of $2^e + \log_2 n$ bits
- One accumulator per exponent:
Malcolm method ... need 2^e accumulators of $f + \log_2 n$ bits

Demmel–Hida method



Consider arithmetic with f -bit mantissa and e -bit exponent ($e = 11$ for fp64).

- One big accumulator: **Kulisch** method
... need one accumulator of $2^e + \log_2 n$ bits
- One accumulator per exponent:
Malcolm method ... need 2^e accumulators of $f + \log_2 n$ bits
- **Demmel–Hida**: general method, balance the number and size of accumulators.

Input: n summands x_i , number of exponent bits m to extract

Output: $y = \sum_{j=1}^{2^m} A_j$

Initialize $A_j = 0$ for $j = 1, \dots, 2^m$

for $i = 1 : n$ **do**

$j \leftarrow m$ leading bits of exponent(x_i)

$A_j \leftarrow A_j + x_i$

end for

With 2^m accumulators, need F -bit mantissa with

$$F \geq f + \lceil \log_2 n \rceil + 2^{e-m} - 1$$

Demmel–Hida method

number of bits					
	signif.	(t)	exp.	range	$u = 2^{-t}$
fp128	quadruple	113	15	$10^{\pm 4932}$	1×10^{-34}
fp64	double	53	11	$10^{\pm 308}$	1×10^{-16}

Numerical example with fp64 and fp128 arithmetics:

- Assume $\log_2 n \leq 29$ ($n \lesssim 0.5 \times 10^9$)
- $F = 113, f = 53, e = 11 \Rightarrow m$ must thus satisfy

$$\begin{aligned} F &\geq f + \lceil \log_2 n \rceil + 2^{e-m} - 1 \\ \Rightarrow 2^{11-m} &\leq 32 \\ \Rightarrow 6 &\leq m \end{aligned}$$

Distillation vs condensation

Distillation methods (AccSum, etc.)

- 😊 Entirely in the working precision
- 😊 Only uses standard arithmetic operations
- 😢 Strongly dependent on the conditioning
- 😢 Limited parallelism

Condensation methods (Demmel–Hida, etc.)

- 😊 Independent on the conditioning
- 😊 High level of parallelism
- 😢 Requires access to the exponent
- 😢 Requires extended precision arithmetic

Distillation vs condensation

Distillation methods (AccSum, etc.)

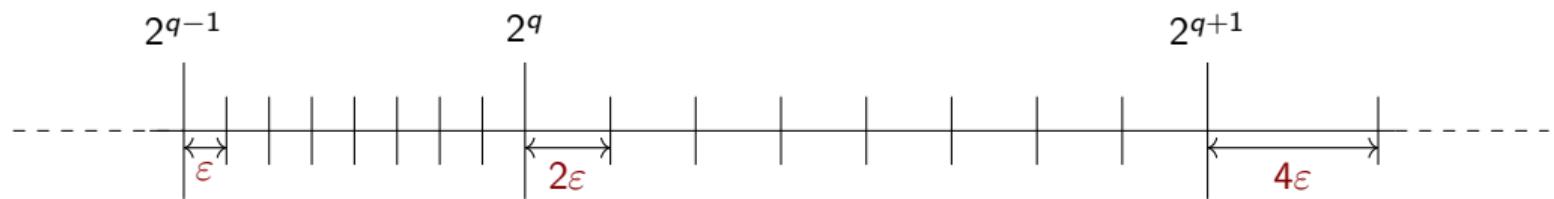
- 😊 Entirely in the working precision
- 😊 Only uses standard arithmetic operations
- 😢 Strongly dependent on the conditioning
- 😢 Limited parallelism

Condensation methods (Demmel–Hida, etc.)

- 😊 Independent on the conditioning
- 😊 High level of parallelism
- 😢 Requires access to the exponent
- 😢 Requires extended precision arithmetic

Can we avoid the use of extended precision arithmetic?

When is $x + y$ exact? Intuition 1



Let $x, y \in \mathbb{F} \cap [2^{q-1}, 2^q]$ such that

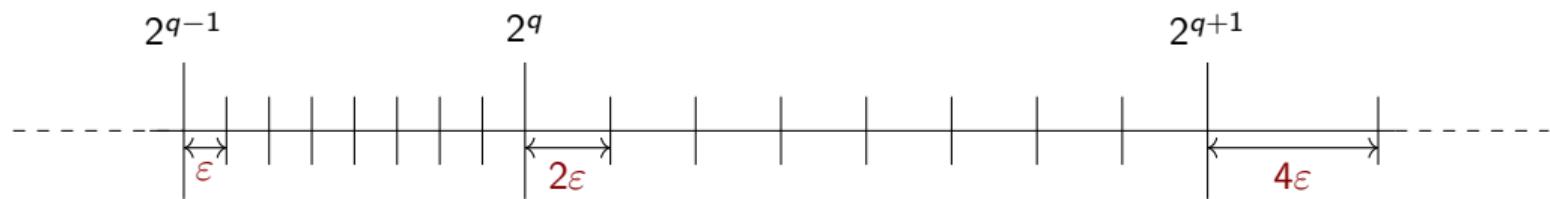
$$x = 2^{q-1} + k_x \varepsilon$$

$$y = 2^{q-1} + k_y \varepsilon$$

Then

$$\begin{aligned} x + y &= 2^{q-1} + k_x \varepsilon + 2^{q-1} + k_y \varepsilon \\ &= 2^q + (k_x + k_y) \varepsilon \in \mathbb{F} \text{ iff } k_x + k_y \equiv 0 \pmod{2} \end{aligned}$$

When is $x + y$ exact? Intuition 1



Similarly if

$$x = 2^{q-1} + k_x \varepsilon$$

$$y = 2^q + k_y 2\varepsilon$$

then $x + y \in \mathbb{F}$ iff

$$\begin{cases} x + y \leq 2^{q+1} \text{ and } k_x \equiv 0 \pmod{2} \\ x + y > 2^{q+1} \text{ and } k_x + 2k_y \equiv 0 \pmod{4} \end{cases}$$

When is $x + y$ exact? Intuition 2

$$2^q \times 10\mathbf{1} + 2^q \times 11\mathbf{1} = 2^q \times 110\mathbf{0} = 2^{q+1} \times 110.\mathbf{0} \in \mathbb{F}$$

$$2^q \times 10\mathbf{1} + 2^q \times 11\mathbf{0} = 2^q \times 101\mathbf{1} = 2^{q+1} \times 101.\mathbf{1} \notin \mathbb{F}$$

$$2^q \times 10\mathbf{1} + 2^{q-1} \times 11\mathbf{1} = 2^{q+1} \times 100.\mathbf{01} \notin \mathbb{F}$$

$$2^q \times 10\mathbf{1} + 2^{q-1} \times 11\mathbf{0} = 2^{q+1} \times 100.\mathbf{00} \in \mathbb{F}$$

When is $x + y$ exact? Theorem

Theorem (Graillat and M.)

Let $x, y \in \mathbb{F}$ of the same sign $\sigma = \pm 1$ such that

$$\begin{aligned}x &= \sigma(\beta^{e_x} + k_x \varepsilon_{e_x}), \\y &= \sigma(\beta^{e_y} + k_y \varepsilon_{e_y}).\end{aligned}$$

Assuming (without loss of generality) that $|x| \leq |y|$, then $x + y \in \mathbb{F}$, and thus the addition is exact, iff one of the following conditions is met:

- (i) $x = 0$;
- (ii) $|x + y| < \beta^{e_y+1}$, $e_y - e_x \leq t - 1$, and $k_x \equiv 0 \pmod{\beta^{e_y - e_x}}$;
- (iii) $|x + y| = \beta^{e_y+1}$, $e_y + 1 \leq e_{\max}$, $e_y - e_x \leq t - 1$, and $k_x \equiv 0 \pmod{\beta^{e_y - e_x}}$;
- (iv) $|x + y| > \beta^{e_y+1}$, $e_y + 1 \leq e_{\max}$, $e_y - e_x \leq t - 2$, and
 $k_x + k_y \beta^{e_y - e_x} \equiv 0 \pmod{\beta^{e_y - e_x + 1}}$.

When is $x + y$ exact? Corollary

$$k_x + k_y \beta^{e_y - e_x} \equiv 0 \pmod{\beta^{e_y - e_x + 1}} \quad \xrightarrow[\beta=2, e_x=e_y]{} \quad k_x + k_y \equiv 0 \pmod{2}$$

Corollary

If $x, y \in \mathbb{F}$ with $\beta = 2$ have the same sign, exponent, and least significant bit, then barring overflow their addition is exact.

Graillat–Mary method

Consider the toy example

$$y = 0.25 + 0.3125 + 0.375 + 0.375 + 0.4375 + 0.4375 + 0.625 + 0.625 + 0.75 + 0.75 + 0.875$$

computed with 3-bit arithmetic:

$$\mathbb{F} = \{0.25, 0.3125, 0.375, 0.4375, 0.5, 0.625, 0.75, 0.875, 1, 1.25, 1.5, 1.75, 2, 2.5, 3\}$$

- LSB=0 $e = 1$
- LSB=1 $e = 0$

- 0.625 $e = -1$
- 0.625
- 0.75
- 0.75
- 0.875

- 0.25 $e = -2$
- 0.3125
- 0.375
- 0.375
- 0.4375
- 0.4375

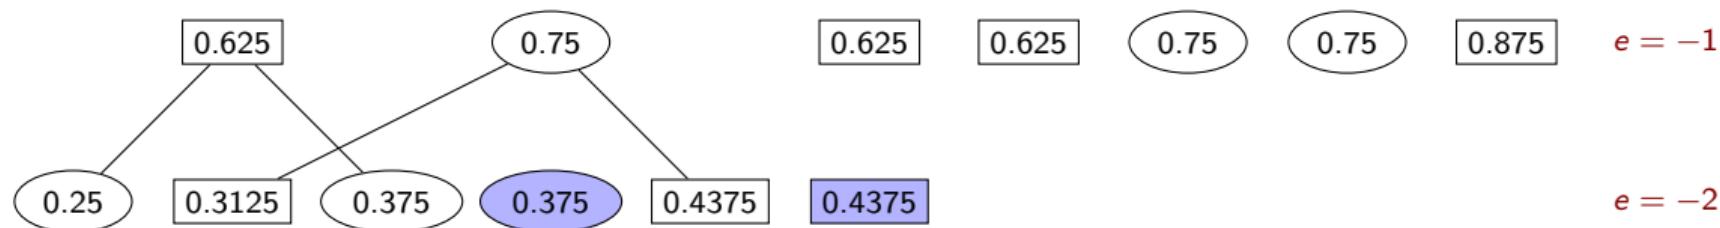
Graillat–Mary method

Consider the toy example

$$y = 0.25 + 0.3125 + 0.375 + 0.375 + 0.4375 + 0.4375 + 0.625 + 0.625 + 0.75 + 0.75 + 0.875$$

computed with 3-bit arithmetic:

$$\mathbb{F} = \{0.25, 0.3125, 0.375, 0.4375, 0.5, 0.625, 0.75, 0.875, 1, 1.25, 1.5, 1.75, 2, 2.5, 3\}$$



Graillat–Mary method

Consider the toy example

$$y = 0.25 + 0.3125 + 0.375 + 0.375 + 0.4375 + 0.4375 + 0.625 + 0.625 + 0.75 + 0.75 + 0.875$$

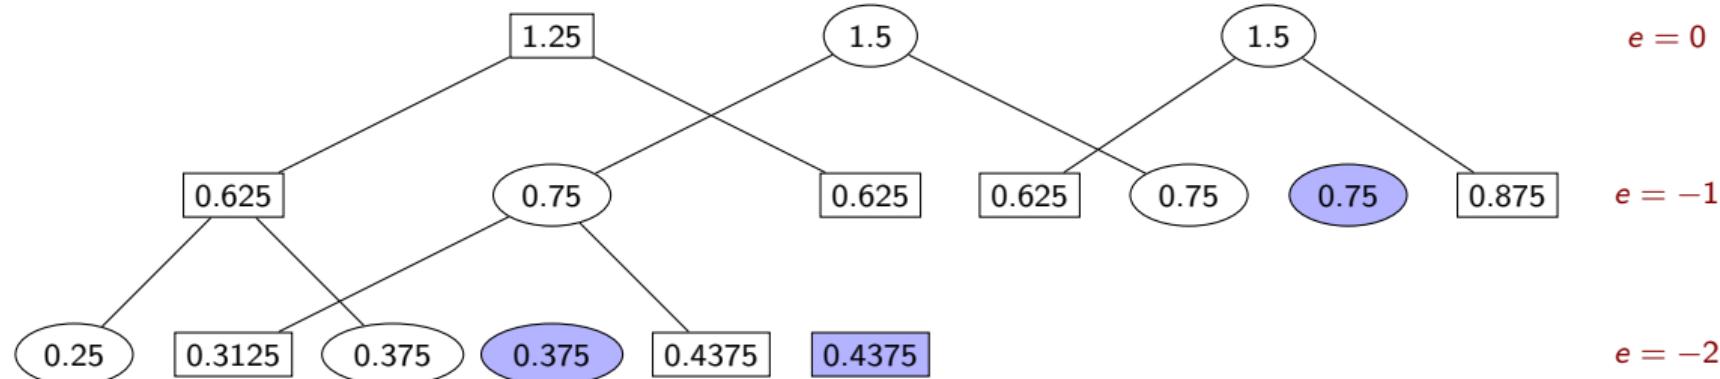
computed with 3-bit arithmetic:

$$\mathbb{F} = \{0.25, 0.3125, 0.375, 0.4375, 0.5, 0.625, 0.75, 0.875, 1, 1.25, 1.5, 1.75, 2, 2.5, 3\}$$

○ LSB=0

□ LSB=1

$e = 1$



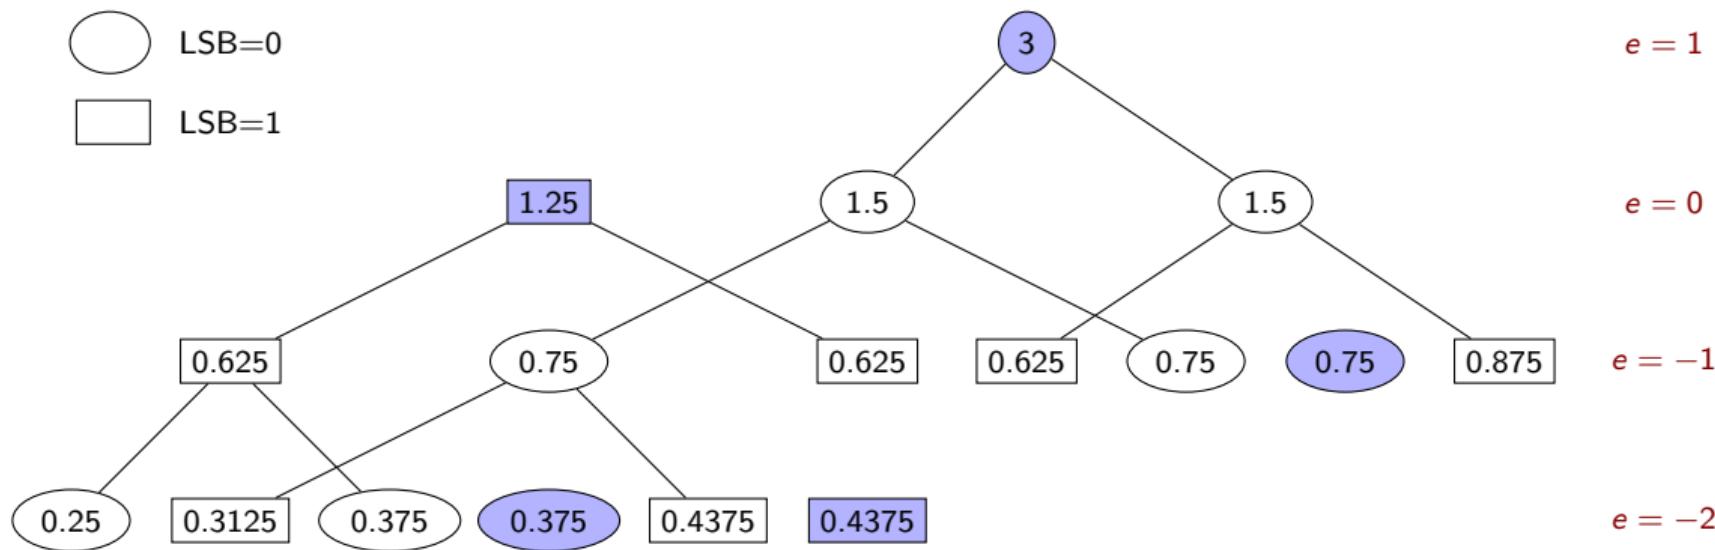
Graillat–Mary method

Consider the toy example

$$y = 0.25 + 0.3125 + 0.375 + 0.375 + 0.4375 + 0.4375 + 0.625 + 0.625 + 0.75 + 0.75 + 0.875$$

computed with 3-bit arithmetic:

$$\mathbb{F} = \{0.25, 0.3125, 0.375, 0.4375, 0.5, 0.625, 0.75, 0.875, 1, 1.25, 1.5, 1.75, 2, 2.5, 3\}$$



Graillat–Mary method

Input: n summands x_i and a distillation method distill

Output: $y = \sum_{i=1}^n x_i$

Initialize $\text{Acc}(e, s, b)$ to 0 for $e = e_{\min} : e_{\max}$,
 $s \in \{-1, 1\}$, $b \in \{0, 1\}$.

for all x_i in any order **do**

$e = \text{exponent}(x_i)$

$s = \text{sign}(x_i)$

$b = \text{LSB}(x_i)$

insert (Acc, x_i, e, s, b)

end for

$x_{\text{condensed}} = \text{gather}(\text{Acc})$

$y = \text{distill}(x_{\text{condensed}})$

```
function insert ( $\text{Acc}, x, e, s, b$ )
if  $\text{Acc}(e, s, b) = 0$  then
     $\text{Acc}(e, s, b) = x$ 
else
     $x' = \text{Acc}(e, s, b) + x$ 
     $\text{Acc}(e, s, b) = 0$ 
     $b' = \text{LSB}(x')$ 
    insert ( $\text{Acc}, x', e + 1, s, b'$ )
end if
end function
```

```
function  $x_{\text{condensed}} = \text{gather}(\text{Acc})$ 
 $i = 0$ 
for all nonzero  $\text{Acc}(e, s, b)$  do
     $i = i + 1$ 
     $x_{\text{condensed}}(i) = \text{Acc}(e, s, b)$ 
end for
end function
```

Conceptual algorithm

$$\mathbb{S} = \{x_1, \dots, x_n\}$$

Repeat for all pairs $(x_i, x_j) \in \mathbb{S}^2$ ($i \neq j$) such that $x_i + x_j$ is exact

$$\mathbb{S} \leftarrow \mathbb{S} \setminus \{x_i, x_j\}$$

$$\mathbb{S} \leftarrow \mathbb{S} \cup \{x_i + x_j\}$$

until no such pair remains

Distill \mathbb{S}

- Can we easily determine when $x_i + x_j$ is exact? YES! It suffices to check the sign, exponent, and LSB of x_i and x_j
- Can we bound the maximum number of leftover summands? YES! At most $4L$ summands where L is the depth of the tree

$$L \leq \lceil \log_2 n \rceil + d$$

where d is independent of n and depends on the range of the values (at most 2047 in binary64)

Distillation vs condensation

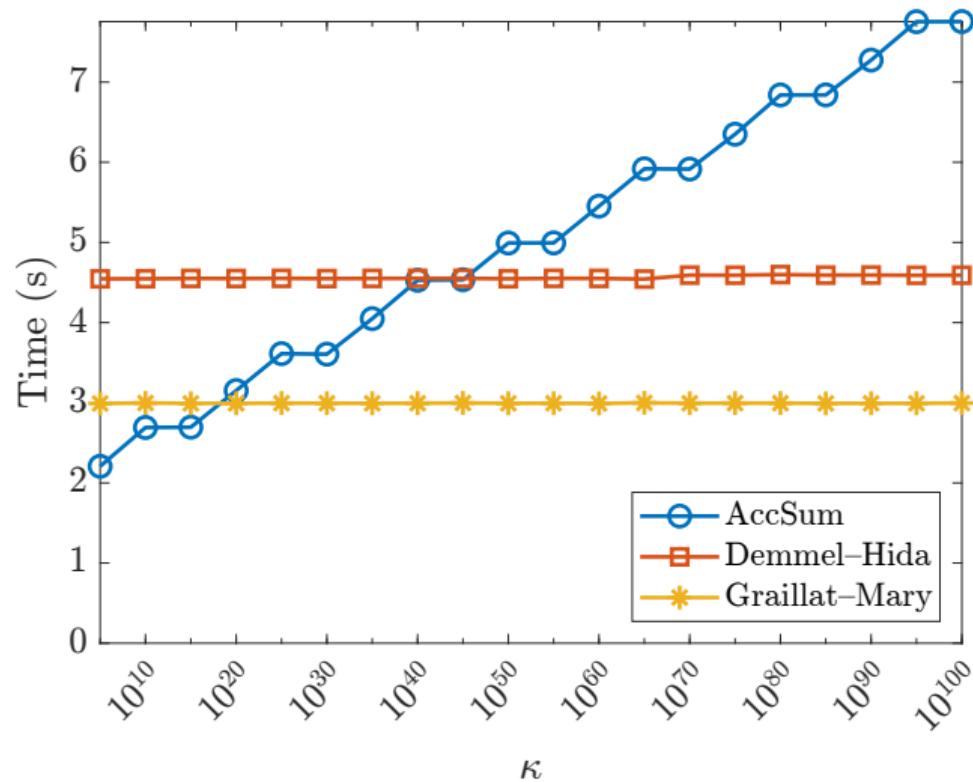
Distillation methods (AccSum, etc.)

- 😊 Entirely in the working precision
- 😊 Only uses standard arithmetic operations
- 😢 Strongly dependent on the conditioning
- 😢 Limited parallelism

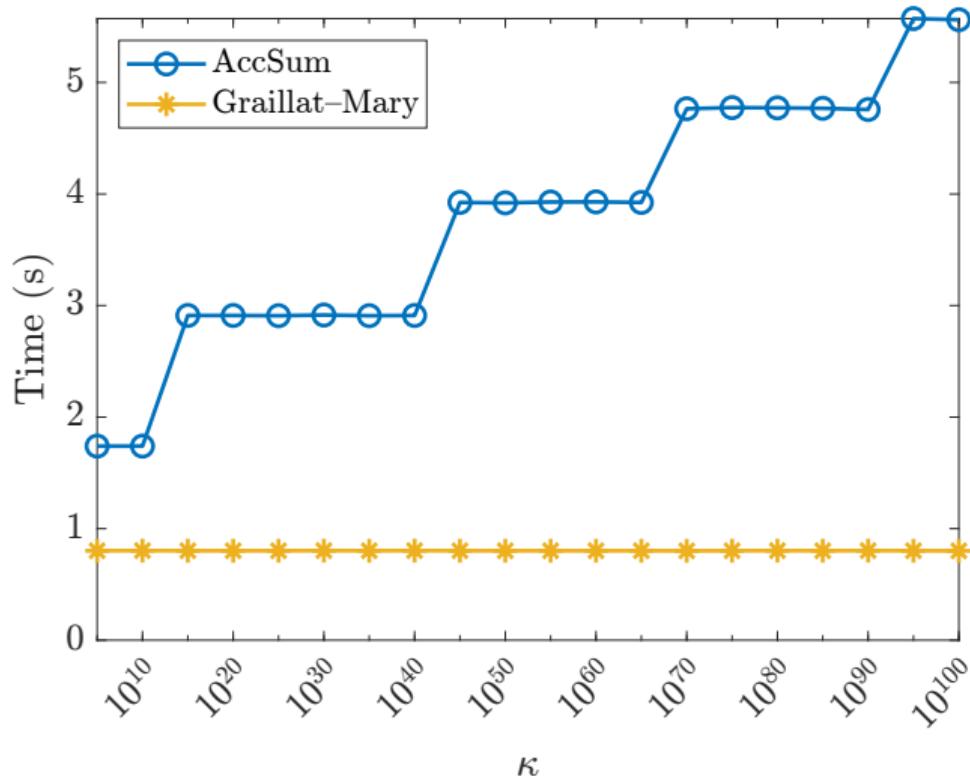
Condensation methods (Demmel–Hida, Graillat–Mary)

- 😊 Independent on the conditioning
- 😊 High level of parallelism
- 😢 Requires access to the exponent + LSB
- 😢 Requires extended precision arithmetic

Performance comparison



Quadruple working precision



Introduction

Dealing with accumulation

Dealing with cancellation

Adaptive precision summation

Conclusion

Adaptive precision algorithms

- Given an algorithm and a prescribed accuracy ε , adaptively select the minimal precision for each instruction depending on the data
- ⇒ **First of all, why should the precisions vary?**

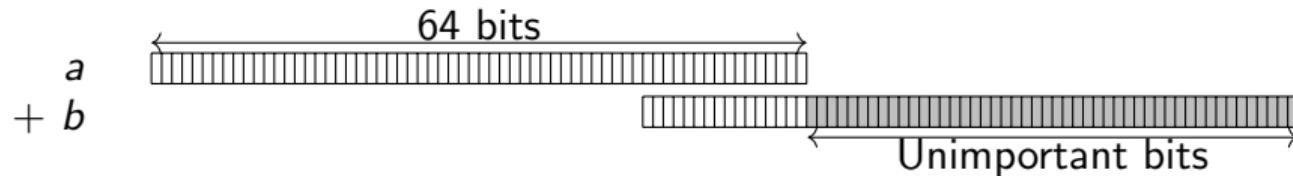
Adaptive precision algorithms

- Given an algorithm and a prescribed accuracy ε , adaptively select the minimal precision for each instruction depending on the data

⇒ **First of all, why should the precisions vary?**

- Because not all computations are equally “important”!

Example:



⇒ **Opportunity for mixed precision:** adapt the precisions to the data at hand by storing and computing “less important” (which usually means smaller) data in lower precision

Sparse matrix–vector product (SpMV)

Goal: compute $y = Ax$, where A is a sparse matrix, with a prescribed accuracy ε

```
for i = 1: m do
     $y_i = \sum_{j \in nnz_i(A)} a_{ij}x_j$ 
end for
```

If computed in precision ε , \hat{y} satisfies

$$|\hat{y}_i - y_i| \leq n_i \varepsilon \sum_{j \in nnz_i(A)} |a_{ij}x_j|$$

and thus

$$\|\hat{y} - y\| \leq c \varepsilon \|A\| \|x\| \quad (c = \max_i n_i)$$

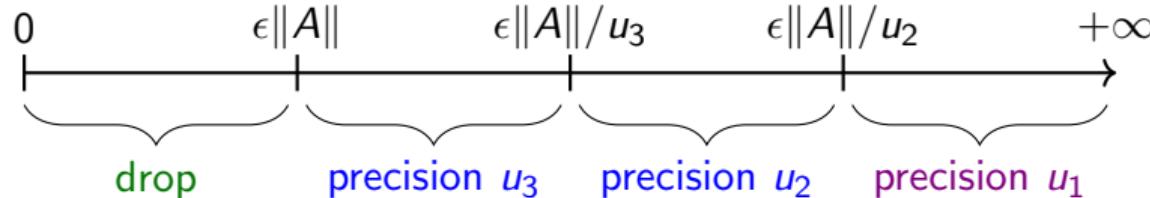
This is a normwise backward error bound: $\hat{y} = (A + E)x$, $\|E\| \leq c \varepsilon \|A\|$.

Adaptive precision SpMV

- Given p available precisions $u_1 < \varepsilon < u_2 < \dots < u_p$, define partition $A = \sum_{k=1}^p A^{(k)}$ where

$$a_{ij}^{(k)} = \begin{cases} \text{fl}_k(a_{ij}) & \text{if } |a_{ij}| \in (\varepsilon \|A\|/u_k, \varepsilon \|A\|/u_{k+1}] \\ 0 & \text{otherwise} \end{cases}$$

⇒ the precision of each element is chosen **inversely proportional to its magnitude**



$$\begin{pmatrix} \times & & \times \\ & \times & \\ \times & & \times \end{pmatrix} = \begin{pmatrix} d & & \\ & d & \\ & & d \end{pmatrix} + \begin{pmatrix} & & s \\ & & s \end{pmatrix} + \begin{pmatrix} & & h \\ & & s \end{pmatrix}$$

Adaptive precision SpMV

```
for i = 1: m do
    for k = 1: p do
         $y_i^{(k)} = \sum_{j \in nnz_i(A^{(k)})} a_{ij}^{(k)} x_j$  in precision  $u_k$ 
    end for
     $y_i = \sum_{k=1}^p y_i^{(k)}$  in precision  $u_1$ 
end for
```

- Compute $y^{(k)} = A^{(k)}x$ in precision u_k . The computed $\hat{y}^{(k)}$ satisfies

$$|\hat{y}_i^{(k)} - y_i^{(k)}| \leq (n_i^{(k)})^2 \varepsilon \|A\| \|x\|$$

- Compute $y = \sum_{k=1}^p y^{(k)}$ in precision u_1 . The computed \hat{y} satisfies

$$\begin{aligned}\hat{y}_i &= \sum_{k=1}^p \hat{y}_i^{(k)} + e_i, \quad |e_i| \leq pu_1 \|A\| \|x\| \\ &= y_i + f_i, \quad |f_i| \leq c\varepsilon \|A\| \|x\|\end{aligned}$$

Custom precision formats

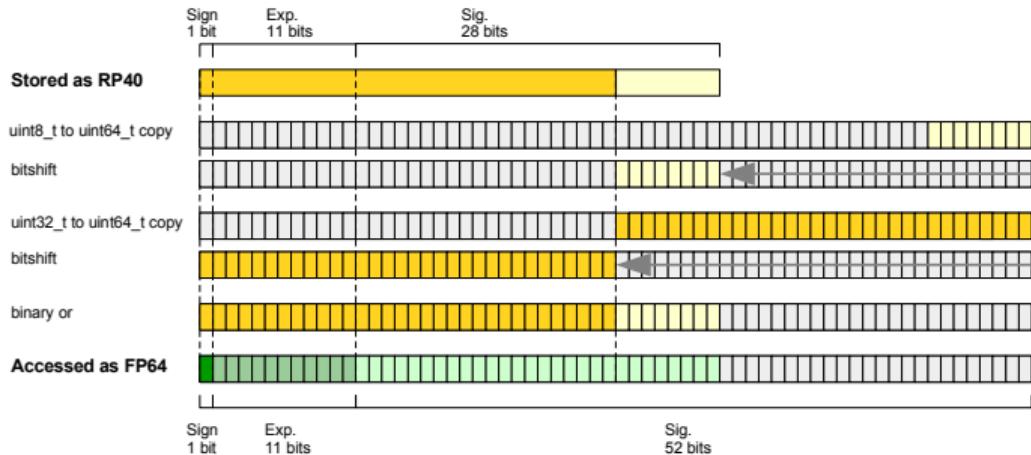
The more precisions we have, the more we can reduce storage \Rightarrow can we exploit custom precision formats?

Emulated formats				
Format	Bits			
	Signif.(t)	Exponent	Range	$u = 2^{-t}$
bf16	8	8	$10^{\pm 38}$	4×10^{-3}
fp24	16	8	$10^{\pm 38}$	2×10^{-5}
fp32	24	8	$10^{\pm 38}$	6×10^{-8}
fp40	29	11	$10^{\pm 308}$	2×10^{-9}
fp48	37	11	$10^{\pm 308}$	8×10^{-12}
fp56	45	11	$10^{\pm 308}$	3×10^{-14}
fp64	53	11	$10^{\pm 308}$	1×10^{-16}

How to efficiently implement custom precision storage?

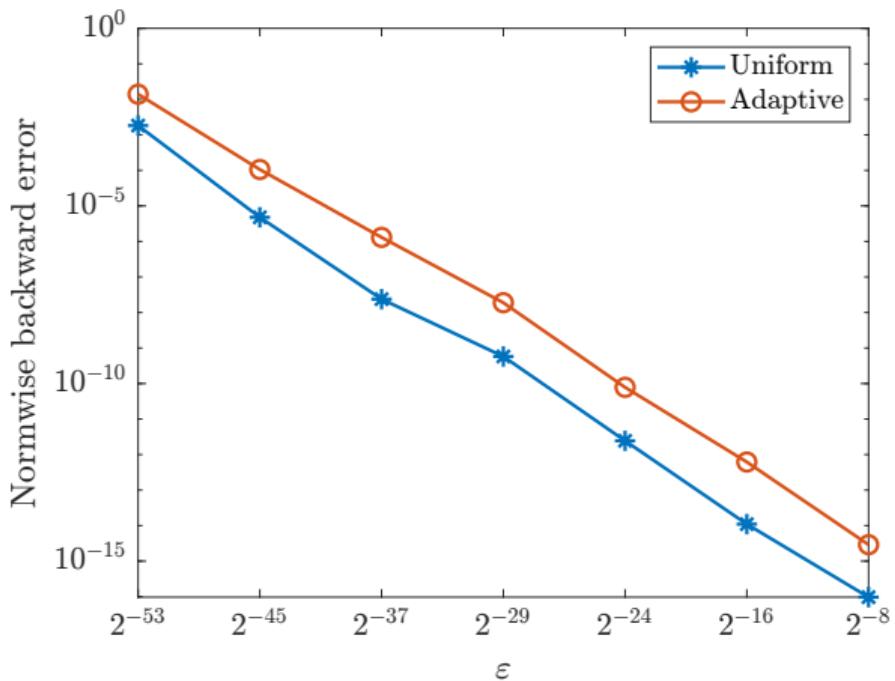
Custom precision accessor

```
union union64 {  
    uint64_t i;  
    double f;  
};  
  
double RpToFp (rp40 rp, size_t i){  
    union union64 u64;  
    uint64_t i64h, i64l;  
    i64h = (uint64_t)rp.i32[i];  
    i64h = i64h << 32;  
    i64l = (uint64_t)rp.i8[i];  
    i64l = i64l << 24;  
    u64.i = i64h | i64l;  
    return u64.f;  
}
```



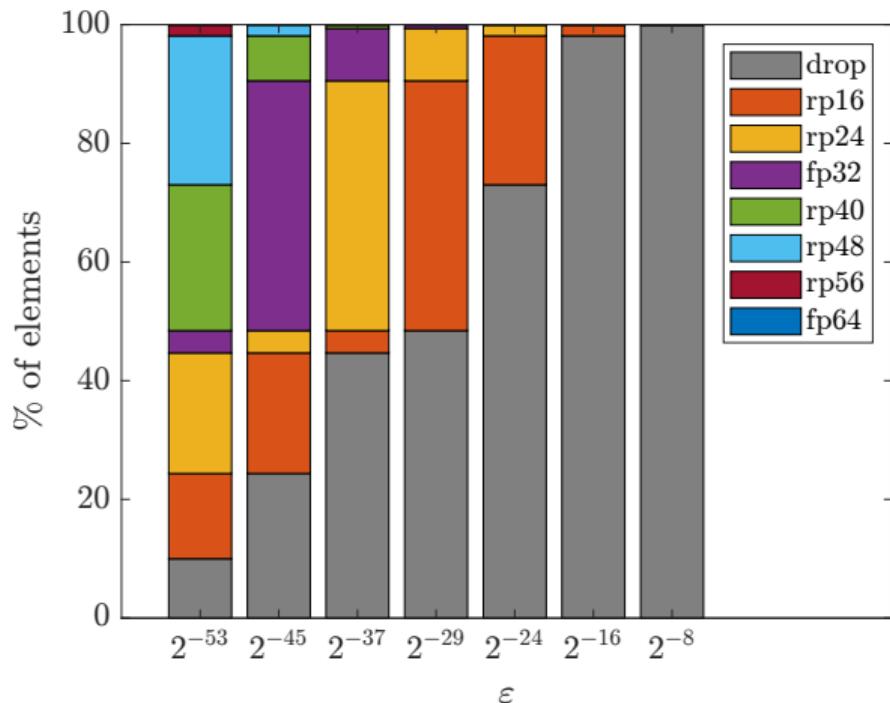
Graillat, Jézéquel, M., Molina, Mukunoki (2024)

Experimental results (Long_Coup_dt6 matrix, $n \approx 1.5M$)



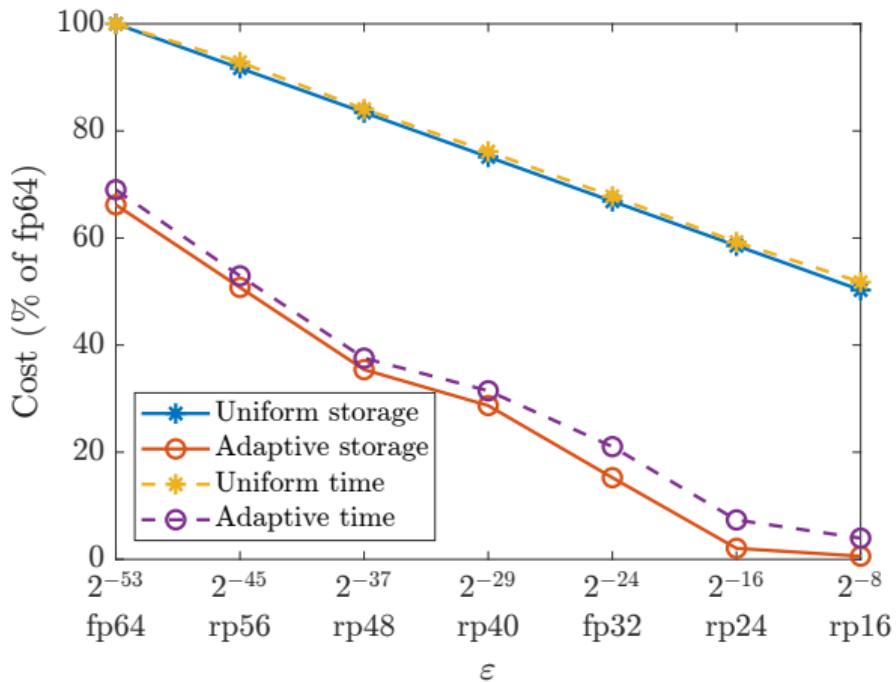
- Controlled accuracy

Experimental results (Long_Coup_dt6 matrix, $n \approx 1.5M$)



- Controlled accuracy

Experimental results (Long_Coup_dt6 matrix, $n \approx 1.5M$)



- Controlled accuracy
- Storage reduced by at least 30% and potentially much more for larger ϵ .
- **Time cost matches storage.**

Introduction

Dealing with accumulation

Dealing with cancellation

Adaptive precision summation

Conclusion

Conclusion

$$\eta_{\text{fwd}} \leq \eta_{\text{bwd}} \kappa, \quad \eta_{\text{bwd}} \leq \gamma_{n-1} = (n-1)u + O(u^2), \quad \kappa = \frac{\sum |x_i|}{|\sum x_i|}$$

- We have seen various summation methods with different properties/objectives: handling error accumulation, cancellation, using mixed precision...
- A common theme has been the **reordering of the summands by grouping them into blocks/buckets**,
 - either fixed-size groups of arbitrary summands
 - or groups of summands of similar magnitude.
- We have seen several possible uses of **mixed precision arithmetic**:
 - Mixed precision blocked summation (FABsum): reduce accumulation
⇒ η_{bwd} independent of n
 - Bucket summation with extended precision (Demmel-Hida): reduce cancellation
⇒ η_{fwd} independent of κ
 - Bucket summation with adaptive precision: exploit lower precisions while controlling η_{bwd}