

An Integer Arithmetic-Based Sparse Linear Solver Using GMRES and Iterative Refinement

Giulia Lionetti Alberto Taddei

AFAE - Floating-Point Arithmetic and Error Analysis

Jan 27, 2026

Paper

Takeshi Iwashita, Kengo Suzuki, Takeshi Fukaya

An Integer Arithmetic-Based Sparse Linear Solver Using a GMRES Method and Iterative Refinement

Roadmap

- ➊ Motivation: Why integer arithmetic?
- ➋ Problem: What breaks without floating-point?
- ➌ Solution: Three-layer architecture
- ➍ Results: Does it actually work?
- ➎ Discussion: Limitations and implications

Is Moore's Law dead?

Dying? We don't know. But hardware is changing.

Three converging trends:

- Energy efficiency gains are slowing
- Novel architectures emerging (SFQ circuits, neuromorphic chips)
- These new technologies may **only support integer arithmetic**
 - Integer circuits are simpler and more energy-efficient
 - FP units are complex and power-hungry

Can we do *real* scientific computing without floating-point?

The problem

Solve $\hat{A}\hat{x} = \hat{b}$ to double precision...
...using only integer arithmetic in the main loop.

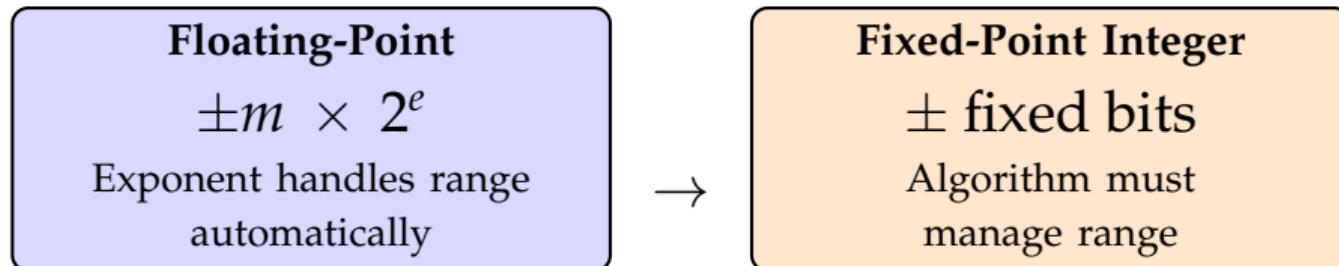
Challenges:

- Fixed range \rightarrow overflow
- Bit shifts \rightarrow lost precision
- No dynamic exponent

Requirements:

- GMRES-like convergence
- Double-precision accuracy
- Robust to ill-conditioning

What we lose without FP



	Floating-Point	Integer/Fixed
Dynamic range	Automatic (exponent)	Manual (shifts)
Overflow risk	Low	High
Energy per op	High (~50 pJ)	Low (~5 pJ)
Hardware complexity	Complex	Simple

This is both the **challenge** and the **opportunity**.

State of the art

Mixed-precision iterative refinement:

- Compute in FP16/FP32, correct in FP64
- Well-studied (Göddeke 2007, Carson & Higham 2018, etc.)
- Works great... but still needs FP hardware

The gap

Nobody's built Krylov solvers that work *without* floating-point kernels.

Three-layer architecture

Layer 1: Iterative Refinement

Compute residual in FP, scale for integer solver, accumulate corrections



Layer 2: Matrix Decomposition

$$A = \bar{A}_0 + 2^{-\alpha_1} \bar{A}_1 + \dots \quad (\text{integers} + \text{power-of-2 weights})$$



Layer 3: int-GMRES

Arnoldi, Givens, matvec \rightarrow all in fixed-point

Integer loop stays bounded; FP recovers accuracy.

Layer 1: Iterative refinement

Standard idea: $x \leftarrow x + \gamma^{(k)} x^{(k)}$ where $x^{(k)}$ solves a scaled residual system.

Why it matters here:

Residual $b' = b - Ax$ shrinks each iteration \rightarrow scale by $1/\gamma = \max |b'_i|$ before sending to integer solver.

Result: integer solver always sees magnitudes near 1.

Refinement = automatic range control

Layer 2: Matrix decomposition

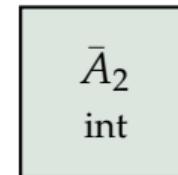
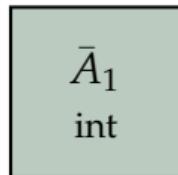
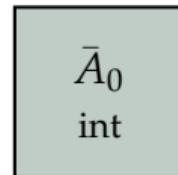
Represent A as:

$$A = \bar{A}_0 + 2^{-\alpha_1} \bar{A}_1 + 2^{-\alpha_2} \bar{A}_2 + \dots$$

$\times 1$

$\times 2^{-\alpha_1}$

$\times 2^{-\alpha_2}$



...

most
significant

fine
detail

- Power-of-2 scaling = cheap bit shifts
- Progressive accuracy: use \bar{A}_0 early, add terms as needed

Layer 3: int-GMRES

Standard GMRES structure (Arnoldi orthogonalization, Givens rotations, least squares).

Key difference: All inner kernels use fixed-point arithmetic.

Integer kernels: matvec, dot products, norms, Givens

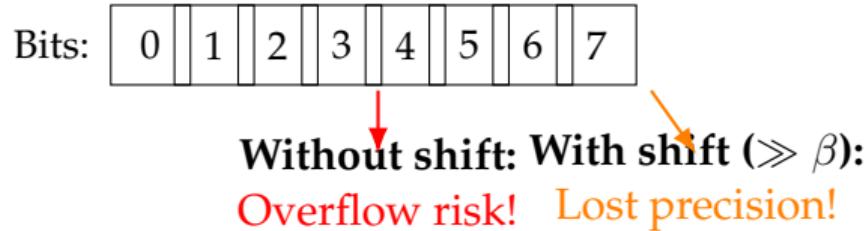
FP only: initial residual, final least squares, solution update

Problem: Dot products and norms overflow easily.

Fixed-point arithmetic

Numbers stored as $Q_{d_m.d_f}$ (sign + d_m integer + d_f fractional bits).

The overflow problem:



$$t_r = ((t_1 \gg \beta_1) \cdot (t_2 \gg \beta_2)) \gg (d_f - \beta_1 - \beta_2)$$

More shift

\Rightarrow safer, loses bits

Less shift

\Rightarrow risky, keeps bits

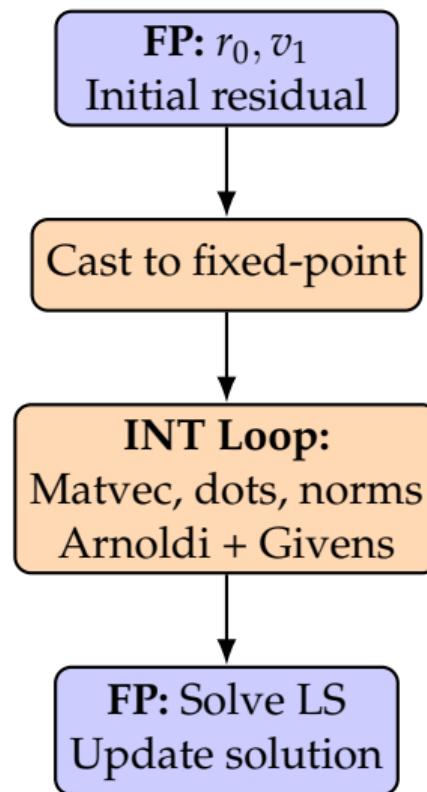
Smart shifting

Exploit GMRES structure:

- Krylov vectors are normalized → many leading zeros → safe
- Givens coefficients satisfy $|c|, |s| \leq 1$ → safe
- Dot products between normalized vectors → mostly safe

Use algorithm invariants to minimize shifts
= maximize effective precision

Core algorithm: Visual overview



Integer kernels:

- Matrix-vector
- Dot products
- Norms
- Givens rotations

Arithmetic stays bounded in the loop; FP ensures accuracy at boundaries.

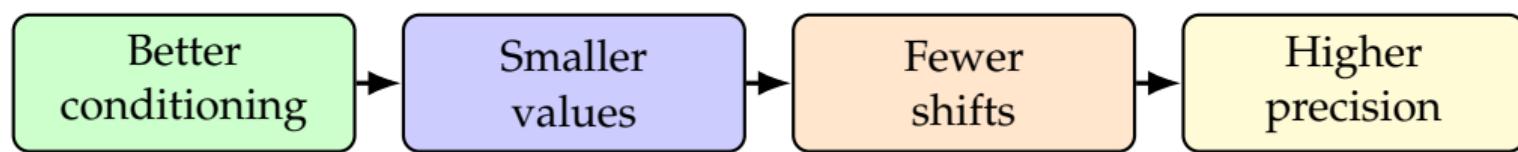
Core algorithm: Pseudocode

```
1:  $r_0 \leftarrow b^{(k)} - A^{(k)}x^{(k)}$ ,  $v_1 \leftarrow r_0 / \|r_0\|$                                 ▷ FP
2:  $\bar{v}_1 \leftarrow \text{cast}(v_1)$                                                  ▷ Convert to fixed-point
3: for  $j = 1, \dots, m$  do
4:    $\bar{w} \leftarrow \bar{A}^{(k)}\bar{v}_j$                                               ▷ INT: matvec
5:   for  $i = 1, \dots, j$  do
6:      $\bar{h}_{i,j} \leftarrow \langle \bar{w}, \bar{v}_i \rangle$                                ▷ INT: dot with shifts
7:      $\bar{w} \leftarrow \bar{w} - \bar{h}_{i,j}\bar{v}_i$                                 ▷ INT: orthogonalize
8:   end for
9:    $\bar{h}_{j+1,j} \leftarrow \|\bar{w}\|$ ,  $\bar{v}_{j+1} \leftarrow \bar{w} / \bar{h}_{j+1,j}$           ▷ INT
10:  Apply Givens rotations                                         ▷ INT
11: end for
12: Solve least squares, update  $x^{(k)}$                                 ▷ FP
```

Preconditioning is critical

Standard reason: faster convergence.

Integer arithmetic reason: reduces overflow risk.



They use **ILU(0)**, implemented entirely in integer arithmetic.

Experimental setup

- 10 sparse matrices from SuiteSparse
- Target: relative residual $< 10^{-8}$ (measured in FP64)
- Fixed-point: $WL = 64, d_f = 30$
- Compare iteration counts: double-GMRES vs int-GMRES

Note

This paper measures **convergence**, not performance.
(Speed/power measurements need target hardware.)

Without preconditioning

Matrix	Double	int-GMRES	Ratio
atmosmodj	2100	2100	1.0×
atmosmodl	420	420	1.0×
wang3	510	630	1.24×
cage14	30	60	2.0×

Mixed results. Some identical, some significantly slower.

Why? Aggressive shifts needed to prevent overflow → accuracy loss → slower convergence.

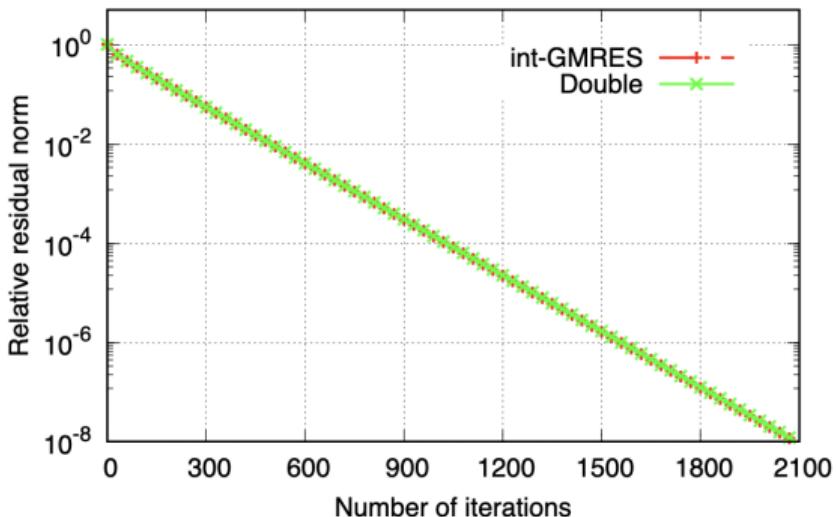
With ILU preconditioning

Matrix	Double+ILU	int+ILU	Ratio
atmosmodj	300	300	1.0×
atmosmodl	120	120	1.0×
wang3	120	120	1.0×
cage14	30	60	2.0×

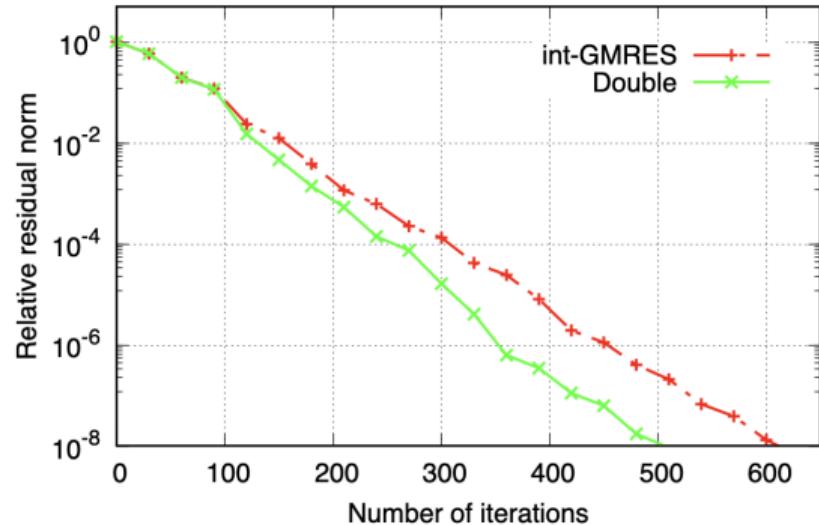
Much better! Most cases identical to double precision.

Preconditioning enables integer arithmetic by stabilizing ranges.

Convergence curves

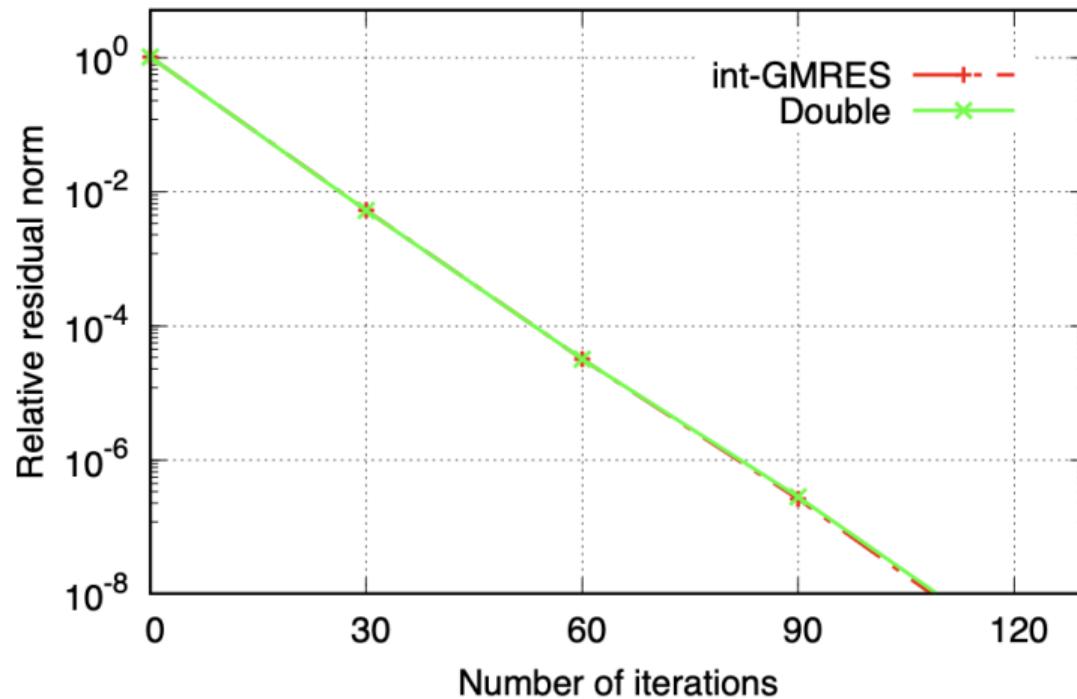


atmosmodj (no precond): identical



wang3 (no precond): slower

With ILU: wang3



With preconditioning, int-GMRES tracks double precision closely.

What they contributed

- ① Working integer-GMRES implementation
- ② Iterative refinement framework for integer arithmetic
- ③ Operation-specific shift strategy exploiting GMRES invariants
- ④ Empirical evidence: ILU preconditioning is essential

Main insight: Precision management moves from hardware to algorithm.

Limitations

- Only tested on moderately conditioned problems
- Several tuning parameters (d_f , shifts, decomposition depth)
- No actual performance/energy measurements
- Theoretical convergence guarantees unclear

But: demonstrates feasibility. Integer Krylov solvers can work.

Conclusions

Integer-only solvers are viable
if you manage range explicitly

- Iterative refinement controls magnitudes
- Smart shifts exploit algorithmic structure
- Preconditioning is crucial

Opens path to ultra-low-power scientific computing.

Thank you!

Questions?

References

-  T. Iwashita, K. Suzuki, T. Fukaya,
An Integer Arithmetic-Based Sparse Linear Solver Using a GMRES Method and Iterative Refinement.
-  Y. Saad,
Iterative Methods for Sparse Linear Systems, SIAM.
-  E. Carson, N. J. Higham,
Iterative refinement in three precisions, SIAM SISC (2018).
-  A. Haidar et al.,
Mixed-precision iterative refinement on GPUs, SC18.