

Floating-point arithmetic and error analysis (AFAE)

Elementary Functions

Stef Graillat

LIP6/PEQUAN – Sorbonne University

Lecture Master 2 CCA



$+, -, \times, /$ and $\sqrt{}$ - it is not all

- Basic operations:

- Addition, subtraction, multiplication, division, square root
- For those operations, we know the exact result:
 - it is a rational number or a zero of a polynomial of degree 2

- it is not all: we also need other functions:

- $\exp(x)$ for bacteria or radioactivity, ...
- $\sin(x)$ for wave, ...
- $\text{atan}(x)$ for reflexion of light, ...
- $\text{erf}(x)$ for statistic and finance, ...

Elementary functions must be computed by a computer. It is implemented in `math.h` and in `libm` library.



How can we compute those elementary functions?

- Question 1: give the first decimal digits of $\sqrt{17}!$
 - $4^2 = 16, \quad 5^2 = 25$
 - $4.5^2 = 20.25, \quad 4.25^2 = 18.0625$
 - $4.125^2 = 17.015625$ so 4.12
- Question 2: give the first decimal digits of $\log(17)!$
 - Not simple! How to do that
 - This is the aim of this lecture.

Outline of the lecture

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

Outline

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

Functions in libm

In **libm mathematical libraries**, we often find:

- elementary functions:

- $\exp, 2^x, 10^x, \log, \log_2, \log_{10}$
- $\sinh, \cosh, \text{asinh}, \text{acosh}$
- $\sin, \cos, \tan, \text{asin}, \text{acos}, \text{atan}$
- x^y

- They are called **elementary** from Liouville.

- Some special functions:

- $\text{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$
- Bessel functions J and Y
- Γ function

- Some composite functions:

- $\exp(x) - 1, \log(1 + x)$
- $\sin(\pi x), \cos(\pi x), \tan(\pi x), \frac{\text{asin}(x)}{\pi} \dots$
- $\text{atan}(\frac{y}{x})$

Framework: floating-point arithmetic

- Floating-point arithmetic is a representation of real numbers
- It is a semi-logarithmic system
 - An exponent E gives the order of magnitude.
 - A significand m indicates the digits of the number.
- Examples :

$$\begin{aligned}105.415 &= 10^3 \cdot 1.05415 = 10^E \cdot m \\0.000000124565 &= 10^{-7} \cdot 1.24565 = 10^E \cdot m \\62 &= 2^5 \cdot 1.11110_2 = 2^E \cdot m\end{aligned}$$

- The representation is written in radix β ; here $\beta = 2$.
- The number of digits is the precision k.
- The set of floating-point numbers with precision k is denoted by \mathbb{F}_k .

IEEE 754 standard

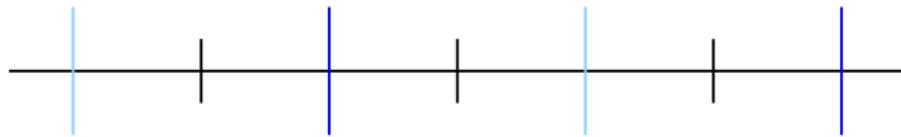
- Computation with floating-point numbers are specified by IEEE 754 standard
- Several formats \mathbb{F}_k are specified

Double precision \mathbb{F}_{53}

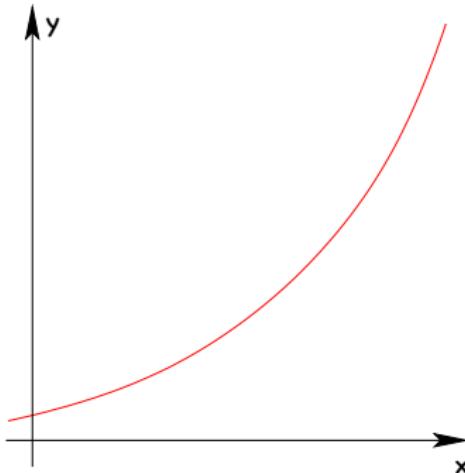


- A rounding in IEEE 754 \circ_k is an application $\mathbb{R} \rightarrow \mathbb{F}_k$:

$$x \mapsto \circ_k(x)$$



An elementary function f on a computer



$$\forall x \in \mathbb{Fk}, \quad x \xrightarrow{f} f(x) \xrightarrow{\circ_k} \circ_k(f(x))$$

Disclaimer

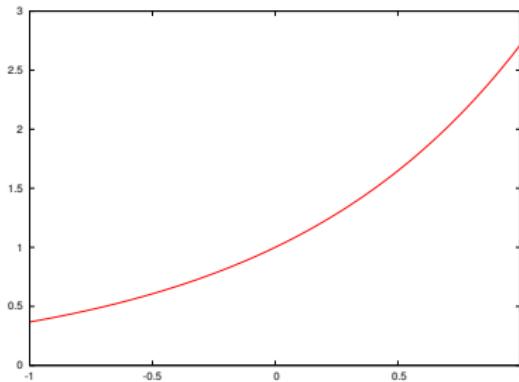
- No multi-precision
 - We know the precision of the input/output format
 - We implement exp for example in binary64 with 53 bits
 - The implementation of exp for k bits in output is more complicated.
 - See [MPFR library](#) for that.
- No special techniques developed for hardware
 - The `libm` are often implemented in software
 - Some processors may have some specific units
 - This is less and less used in the general case
 - When it is used, it is Intel/AMD and it is like in software
 - There are special algorithms for hardware ([CORDIC](#))

Outline

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

Polynomial approximation

- How to implement for example $f = \exp \dots$
- ... on a restricted domain first, for example $I = [-1; 1]$?



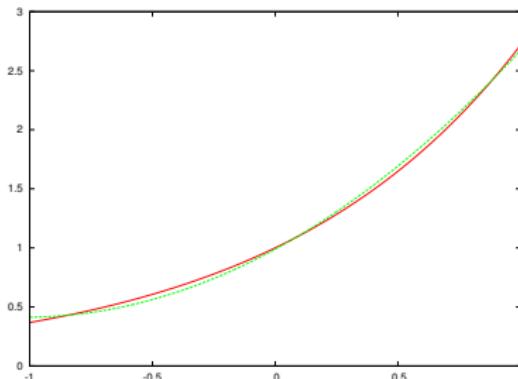
- Solution: replace f by an approximation polynomial p
- Choice for polynomials:
 - Weierstrass theorem tells us it always exists
 - Taylor polynomials
 - Polynomials minimizing the maximal error on the domain

Why choosing polynomials?

- Why choosing polynomials and not other types of approximations?
 - Why not continued fractions (truncated)?
 - Why not approximation with other basis functions?
- An answer (only) purely technological
 - Fast hardware +, \times and FMA
 - Less performance for division (19 cycles on Intel Core i5/i7)
 - Possibility to compute the rounding error for + and \times (and $/$) ... but not for sin and log
- An answer less categorical
 - Some functions are not well approximated by polynomials:
asin
 - We talk about software, but on hardware things can be different.
 - The toolchain is not as comprehensive for other approaches.

Need for an argument reduction

- Need to implement f on the whole range of floating-point numbers
- In binary64, $f = \exp$ is defined on $I = [-744.5; 709]$
- The polynomial approximation is not sufficient:



- When the domain is large, approximation error becomes very large for a given degree
- otherwise: we need to increase the degree to become very large

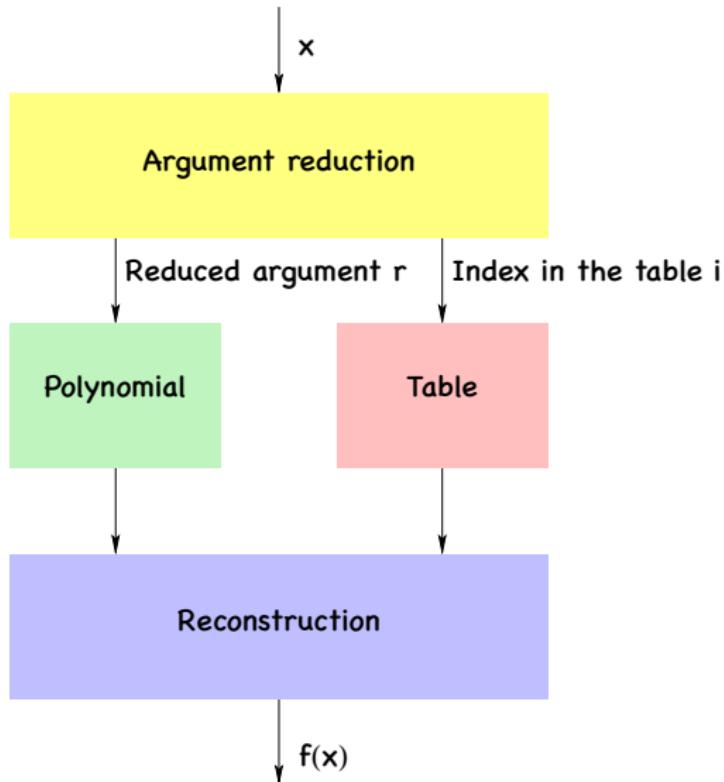
Reduction and tabulation

- Argument reduction brings everything into a small domain
 - Algebraic properties of the function
 - Periodicity: sin
 - Self-similarity: exp
 - Symmetry: asin
 - Semi-logarithmic feature of floating-point format
 - Somehow, convertToInteger and exp are very similar
 - If nothing else: split the domain
- Example for exp:

$$e^x = 2^{\frac{x}{\log 2}} = 2^{\left\lfloor \frac{x}{\log 2} \right\rfloor} \cdot 2^{\frac{x}{\log 2} - \left\lfloor \frac{x}{\log 2} \right\rfloor} = 2^E \cdot e^{x-E \log 2} = 2^E \cdot e^r$$

- The reduction is often linked to a tabulation
 - Tabulation: pre-computation of a function g at some discrete points
 - Reduction: the polynomial p must be valid only between those discrete points

Scheme for the computation of a function



A toy exponential function

```
// A very crude "toy" implementation of exp(x)
//
// About 45 bits of accuracy. No checks for NaN, Inf whatsoever.
//
double Exp(double x) {
    double z, n, t, r, P, tbl, y;
    uint32_t E, idx, N;
    doubleCaster shiftedN, twoE;

    // Argument reduction
    z = x * TWO_4_RCP_LN_2;           // z = x * 2^4 * 1/ln(2)
    shiftedN.d = z + TWO_52_P_51;     // shiftedN.d = nearestint(z) + 2^52 + 2^51
    n = shiftedN.d <> TWO_52_P_51;   // n = nearestint(z) as double
    N = shiftedN.i[LO];              // N = nearestint(z) as integer
    E = N >> 4;                     // E = floor(2^4 * N)
    idx = N & 0x0F;                  // idx = N <> E * 2^4
    t = n * TWO_M_4_LN_2;            // t = n * 2^4 * ln(2)
    r = x <> t;                    // r = x <> t

    // Polynomial approximation          p(r) approximates exp(r)
    P = c0 + r * (c1 + r * (c2 + r * (c3 + r * (c4 + r * c5))));

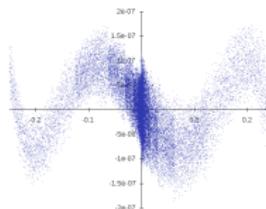
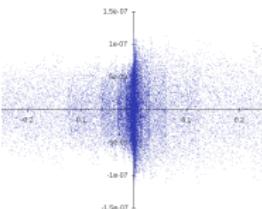
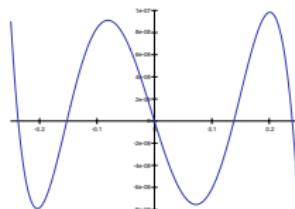
    // Table access
    tbl = table[idx];                // tbl = 2^(2^4 * idx)

    // Reconstruction
    twoE.i[HI] = (E + 1023) << 20;
    twoE.i[LO] = 0;                  // twoE.d = 2^E
    y = twoE.d * (tbl * P);         // y = 2^E * tbl * P

    return y;
}
```

Sources of errors

- Approximations with floating-point numbers, rounding errors



- Five principal sources of errors:
 - Error in the argument reduction
 - Error in the entries of the table
 - Approximation error $\|p/f - 1\|_\infty$
 - Error in the polynomial evaluation $|P(x)/p(x) - 1|$
 - Error in the reconstruction
- Combination of those errors to get the global error

Outline

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

Problem of polynomial approximation

- We have a function $f : \mathbb{R} \rightarrow \mathbb{R}$...
- on a small domain $I = [a; b] \subset \mathbb{R}$.
- We are given a targeted error $\bar{\varepsilon}$ we do not want to exceed.
- We are looking for a polynomial p of degree n minimal such that the error $\|p - f\|_\infty$ is bounded (in absolute value) by $\bar{\varepsilon}$.

Example:

- Function: $f = \exp$
- Domain: $I = [-\frac{1}{2}; \frac{1}{2}]$
- Targeted error: $\bar{\varepsilon} = 2^{-5}$ (5 bits of accuracy)
- $p(x) = 1 + x + \frac{1}{2}x^2$ $n = 2$
- $\|p - f\|_\infty \approx 3.05 \cdot 10^{-2} < 2^{-5}$

Problem of polynomial approximation - 2

- In other words, we have to
 - fix a degree n we hope to be sufficient
 - compute some coefficients c_i of a polynomial

$$p(x) = \sum_{i=0}^n c_i x^i$$

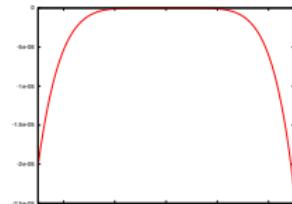
- look at the error $\frac{p}{f} - 1$
- re-begin with a larger n if the error is not small enough

How can we compute the coefficients c_i of the polynomial p ?

Taylor polynomials

- First idea:

- Let us take $x_0 \in I$
- $f(x_0)$ approximates $f(x)$ on I
- $f(x_0) + (x - x_0) \cdot f'(x_0)$ is better
- $f(x_0) + (x - x_0) \cdot f'(x_0) + \frac{1}{2} \cdot (x - x_0)^2 \cdot f''(x_0)$ is even better...



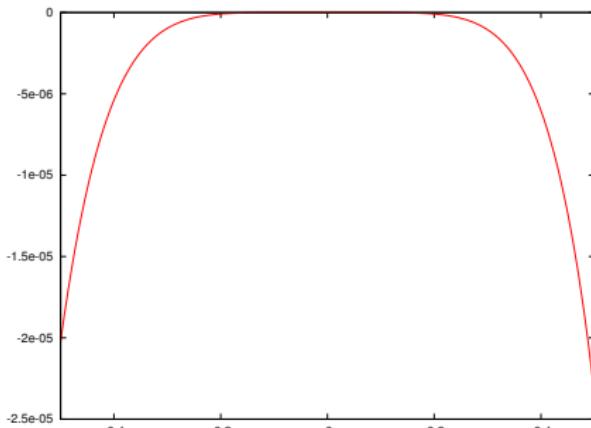
- This is the idea of Taylor polynomials at x_0 :

$$f(x) = \underbrace{\sum_{i=0}^n \frac{f^{(i)}(x_0)}{i!} \cdot (x - x_0)^i}_{=: p(x)} + \underbrace{\frac{f^{(n+1)}(\xi)}{(n+1)!} \cdot (x - x_0)^{n+1}}_{\text{Lagrange remainder resp. Error}}$$

- Functions are defined by simple differential equations:
 - Taylor polynomials known for the functions in a libm library

Problems with Taylor polynomials

Error of Taylor polynomial p of degree 5 with respect to
 $f = \exp$:



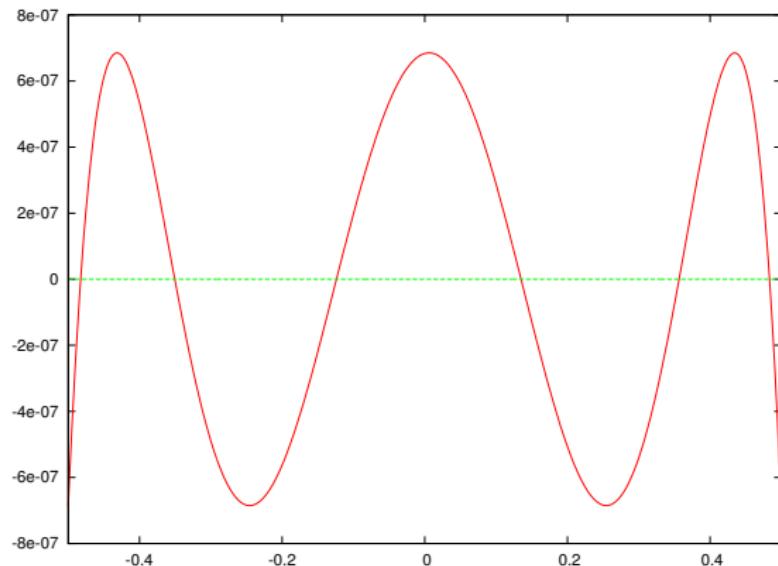
- Taylor polynomial approximates the function **the best at the point x_0** : indeed the error vanishes at x_0
- The error blow up at the boundary of the domain I
- \Rightarrow it would be necessary for p to approximate f in the whole domain I, or at least **the error vanishes several times**

Interpolation polynomial

- A polynomial of degree n
 - $n + 1$ coefficients
 - $n + 1$ degree of freedom
 - defined by $n + 1$ points $(x_j, p(x_j))$
- In order for the error $p(x) - f(x)$ to vanish at points x_j , it is sufficient to have $p(x_j) = f(x_j)$.
- The polynomial p interpolates the function f at x_j
- An interpolation polynomial of f is defined by $n + 1$ values x_j
 - One value by degree of freedom
 - The ordinates $f(x_j)$ are clear for f and x_j .

Example of interpolation

Error of interpolation polynomial p with respect to $f = \exp$:



- The error vanishes at the interpolation points
- By choosing the points, we can constrain the error to oscillate

Computation of interpolation polynomial

How can we compute an interpolation polynomial?

- We have $p(x_j) = \sum_{i=0}^n c_i x_j^i = f(x_j)$ for $n+1$ points x_j
- In other words, we have

$$\underbrace{\begin{pmatrix} 1 & x_1 & x_1^2 & \cdots & \cdots & x_1^n \\ 1 & x_2 & \ddots & & & \vdots \\ \vdots & \vdots & & x_j^i & & \vdots \\ \vdots & \vdots & & \ddots & & \vdots \\ 1 & x_{n+1} & \cdots & \cdots & \cdots & x_{n+1}^n \end{pmatrix}}_{=:A} \cdot \underbrace{\begin{pmatrix} c_0 \\ c_1 \\ \vdots \\ c_n \end{pmatrix}}_{=:p} = \underbrace{\begin{pmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ \vdots \\ f(x_{n+1}) \end{pmatrix}}_{=:f}$$

- We know the **Vandermonde** matrix A and the vector \vec{f} .
 - the solution of a linear system gives the coefficients of p
 - Moreover: there are special algorithms for such kind of linear systems

Going around in circles ...

- We know the RHS \vec{f} .
- So we can compute $f(x_j)$ for all x_j .
- Strange? we were implementing a function to compute f !
- This is a problem:
 - We first need to write a code to evaluate $f(x_j)$
 - Then, we can construct the interpolation polynomials in order to evaluate f everywhere.
- In practice:
 - We write multi-precision procedures for functions
 - Based on Taylor polynomials which are inefficient
 - Then, we use those codes to develop a libm

Which points to interpolate?

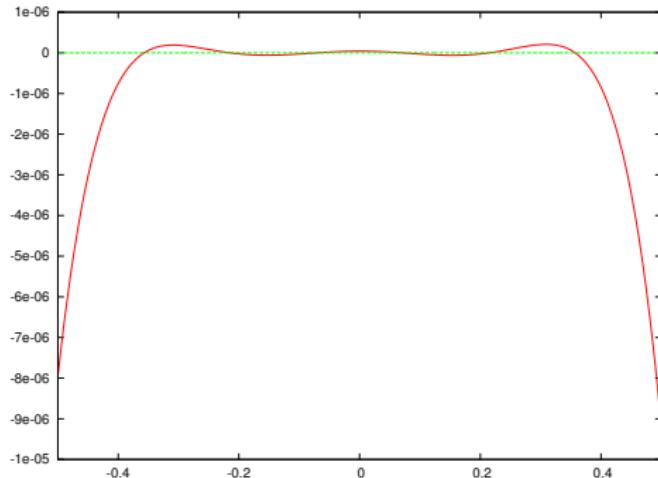
- The choice of points x_j has an influence on the error:
Let p be an interpolation polynomial of f at points x_j .
Then we have

$$f(x) = p(x) + \frac{1}{(n+1)!} f^{(n+1)}(\xi) \prod_{j=0}^n (x - x_j)$$

- The equidistant points are not optimal
- Tchebychev tried to minimize the error
- Remez finally provided an iterative algorithm
 - The algorithm provides the polynomial with minimal error with maximum norm
 - It only choose the best interpolation points

Equidistant points

Error of interpolation p with $f = \exp$ at equidistant points:

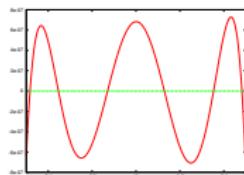


- It is better than Taylor: error 2.5 times lower
- Error oscillates
- Error still blows up at the boundary
⇒ We want to put more points near the boundary

Tchebychev points

- Interpolation error to minimize with $\|\cdot\|_\infty$ norm:

$$f(x) - p(x) = \frac{1}{(n+1)!} f^{(n+1)}(\xi) \prod_{j=0}^n (x - x_j)$$



- P. L. Tchebychev: try to minimize

$$\left\| \prod_{j=0}^n (x - x_j) \right\|_\infty$$

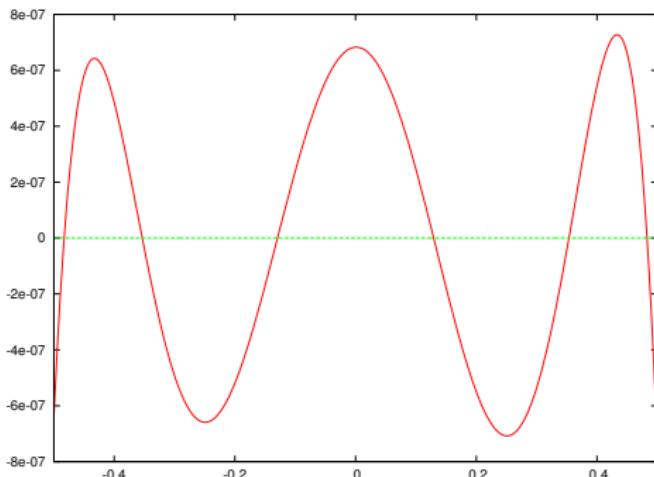
- Points for which it is the case in $I = [a; b]$:

$$x_j = a + \frac{b-a}{2} \cdot \left(\cos \left(\frac{2j-1}{2(n+1)} \pi \right) + 1 \right)$$

- Those are Tchebychev interpolation points.

Tchebychev points - example

Error of interpolation polynomial p with $f = \exp$ at Tchebychev points:



- It is 10 times better than for equidistant points!
- It is not optimal yet: cf. picture
- There are functions for which the gap is very large

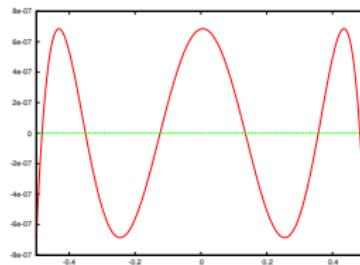
Remez polynomial

- Theorem (Tchebychev/La Vallée-Poussin):

The approximation is optimal iff all the extrema of the error are at the same height.

- E. Ya. Remez:

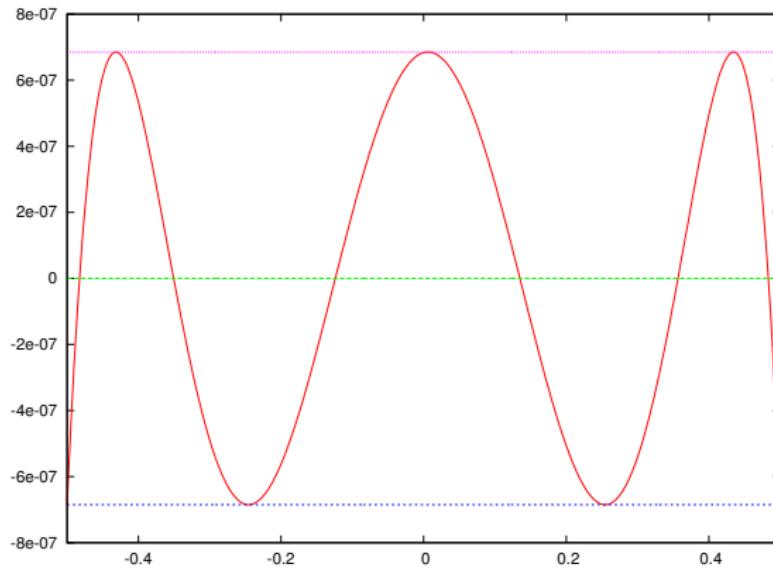
- Interpolate directly $f(x) + (-1)^j \cdot \varepsilon$ where ε is the height of the extrema we are looking for
- While the extrema of the error are not at the same height, exchange interpolation points with points where the extrema are reached.



- This is **Remez algorithm**
- It converges toward the polynomial with **minimal error** measured with **max norm**
- That is why it is also called **minimax algorithm**

Example: Remez polynomial

Error of the Remez polynomial p of degree 5 with respect to



$f = \exp:$

- The error is a little bit smaller than for Tchebychev
- All the extrema are at the same height

Still one problem → fpminimax

- Taylor, Tchebychev, Remez: polynomials with real coefficients

$$p(x) = \sum_{i=0}^n c_i x^i, \quad c_i \in \mathbb{R}$$

- On a computer, we only have floating-point numbers to store the c_i .
- If we only round numbers coefficient by coefficient, $\tilde{c}_i = \circ_{c_i}$, we destroy all the work of Remez algorithm.
 - The oscillations disappear.
 - The error blows up.
- It is important to find polynomials with floating-point coefficients

$$p(x) = \sum_{i=0}^n c_i x^i, \quad c_i \in \mathbb{Fk}$$

- This is a discrete optimization problem
- Since 2006, there is a heuristic method based on lattice reduction.

Outline

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

Polynomials are not all ...

- The degree of the polynomials blows up when the size of the domain I increases.
- Idea:
 - For a periodic function like \sin , it is not necessary to approximate it on $[2\pi; 4\pi]$ when we already know it on $[0; 2\pi]$.
 - A function like 2^x should benefit for the fact that we use a semi-logarithmic representation $2^E \cdot m$.
- Argument reduction: use
 - periodicity,
 - self-similarity,
 - symmetry
 - and the semi-logarithmic property of floating-point numbersto make the domain where the function has to be approximated as small as possible.

Increase of the degree - example

To approximate e^x by a polynomial with 24 bits of accuracy,
we need

on	a degree n of
$[-\frac{1}{4}; \frac{1}{4}]$	5
$[-\frac{1}{2}; \frac{1}{2}]$	6
$[-1; 1]$	8
$[-2; 2]$	11
$[-16; 16]$	42
$[-32; 32]$	78
$[-256; 256]$	> 93

In double precision, we need to provide 53 bits of accuracy
and cover the domain $[-745; 709] \dots$

Argument reduction

- Difficult to precisely define what argument reduction is.
- We try our best:

Definition (Argument reduction)

Given $f : \mathbb{F} \rightarrow \mathbb{F}$, an argument reduction consists in

- a **reduction function** $r : \mathbb{F} \rightarrow \mathbb{F}^n$ which is **simple** to calculate,
- a **reduced function** $g : \mathbb{F}^n \rightarrow \mathbb{F}^m$ whose projections can be computed easily either by **polynomial approximation** or by tabulation,
- and a **reconstruction function** $c : \mathbb{F}^m \rightarrow \mathbb{F}$ which can be computed easily.

These functions are such that

$$f(x) = c(g(r(x))) \quad \forall x \in \mathbb{F}.$$

Example 1: exp without table

We want to compute e^x . We have:

$$\begin{aligned} e^x &= 2^{\log_2(e) \cdot x} \\ &= 2^E \cdot 2^{\log_2(e) \cdot x - E} \\ &= 2^E \cdot 2^{\log_2(e) \cdot \left(x - \frac{1}{\log_2(e)} \cdot E\right)} \\ &= 2^E \cdot e^{x - \frac{1}{\log_2(e)} \cdot E} \\ &= 2^E \cdot e^r \end{aligned}$$

with $E = \lfloor \log_2(e) \cdot x \rfloor$ and $r = x - \frac{1}{\log_2(e)} \cdot E.$

$\underbrace{\hspace{10em}}$
reduction function r

- **r is bounded** by $|r| \leq \frac{1}{2 \cdot \log_2(e)} \approx 0.35.$
- So **the function $g(r) = e^r$ is well approximated** by a polynomial.
- The reconstruction c is simple, one has to multiply e^r by $2^E.$

Tabulation

- Just with multiplications by 2^E and polynomial, it is difficult.
- Instead of computation, we could memorize.
- We try to couple an argument reduction with the read in precomputed tables.
- Those tables are modelized by discrete reduced functions with hundred to thousand possible entries (otherwise, it is too expensive).

Example 2: exp with table

We still want to compute e^x . We have:

$$\begin{aligned} e^x &= 2^E \cdot 2^{k \cdot 2^{-w} - E} \cdot e^{x - \frac{1}{\log_2(e)} \cdot k \cdot 2^{-w}} \\ &= 2^E \cdot 2^i \cdot e^r \end{aligned}$$

with

$$k = \lfloor \log_2(e) \cdot x \cdot 2^w \rfloor$$

$$E = \lfloor k \cdot 2^{-w} \rfloor$$

$$i = k \cdot 2^{-w} - E$$

$$r = x - \frac{1}{\log_2(e)} \cdot k \cdot 2^{-w}$$

where w is the number of bits used for the index of the table.

- r is bounded by $|r| \leq 2^{-w} \cdot \frac{1}{2 \cdot \log_2(e)}$.
- We read 2^i in a precomputed table.

How to find a good argument reduction?

- An argument reduction
 - needs the use of the properties of the function f
 - and of floating-point arithmetic
 - needs several functions which are easily computable or that can be put into tables
- Consequences: find a good argument reduction is difficult
- For that, we can use:
 - The Handbook of Mathematical Functions written in the 1950s by Abramowitz and Stegun
 - experiences and intuitions.

Relations Between Circular Functions

$$4.3.10 \quad \sin^2 z + \cos^2 z = 1$$

$$4.3.11 \quad \sec^2 z - \tan^2 z = 1$$

$$4.3.12 \quad \csc^2 z - \cot^2 z = 1$$

Negative Angle Formulas

$$4.3.13 \quad \sin(-z) = -\sin z$$

$$4.3.14 \quad \cos(-z) = \cos z$$

$$4.3.15 \quad \tan(-z) = -\tan z$$

Addition Formulas

$$4.3.16 \quad \sin(z_1 + z_2) = \sin z_1 \cos z_2 + \cos z_1 \sin z_2$$

$$4.3.29 \quad \sin 4z = 8 \cos^3 z \sin z - 4 \cos z \sin z$$

$$4.3.30 \quad \cos 4z = 8 \cos^4 z - 8 \cos^2 z + 1$$

Products of Sines and Cosines

$$4.3.31 \quad 2 \sin z_1 \sin z_2 = \cos(z_1 - z_2) - \cos(z_1 + z_2)$$

$$4.3.32 \quad 2 \cos z_1 \cos z_2 = \cos(z_1 - z_2) + \cos(z_1 + z_2)$$

$$4.3.33 \quad 2 \sin z_1 \cos z_2 = \sin(z_1 - z_2) + \sin(z_1 + z_2)$$

Addition and Subtraction of Two Circular Functions

$$4.3.34$$

$$\sin z_1 + \sin z_2 = 2 \sin\left(\frac{z_1 + z_2}{2}\right) \cos\left(\frac{z_1 - z_2}{2}\right)$$

Outline

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

Performances

- On typical Intel/AMD or IBM processors
 - a floating-point addition or multiplication takes 4 cycles
 - a division 20 cycles
 - a function call 15 cycles
 - a function like `exp` takes 40 cycles
- Within 40 cycles, one has to do lots of things:
 - an argument reduction
 - one or two read in a table
 - one polynomial evaluation
 - a reconstruction
 - dealing with NaN, $\pm\infty$, etc.
- Consequences:
 - We need to highly parallelize the code
 - The functions `libm` cannot call other functions
 - A clever manipulation of floating-point numbers

Example 1: computation of $\lfloor x \rfloor$

- We want to compute $\lfloor x \rfloor$:
 - that is to say the nearest integer to the floating-point number x .
 - We could use `rint(x)` or `floor(x + 0.5)`
 - It would just cost 30 cycles for the call
- Or we can notice that:
 - $\lfloor x \rfloor$ is the rounding to nearest of a quantity for which $\text{ulp} = 1$
 - For x sufficiently small, $2^{52} + 2^{51} + x$ is such a quantity (for double precision)
 - It is just then needed to remove the shift $2^{52} + 2^{51}$.
- So we do:

```
double x, shifter, tmp, nearest;  
x = ....                                // computation that gives x  
shifter = 6755399441055744.0;           // shifter =  $2^{52} + 2^{51}$   
tmp = x + shifter;                      // tmp = round( $2^{52} + 2^{51} + x$ )  
nearest = tmp - shifter;                 // nearest = nearestint(x) - Sterbe
```

Even better

- For the computation $[x]$ in the previous slide
 - the input x is a double floating-point number
 - the output nearest is a double floating-point number
- Often, we want the output in a variable n of type int
 - We could do $n = (\text{int}) \text{nearest};$ in C
 - The compiler would insert a conversion instruction or a function call
 - \Rightarrow it would take plenty of cycles
- But in fact, with IEEE 754-2008 standard,
 - the unit on the last place is stored at the end in memory
 - before, there is the sign, the exponent and the other bits of the significand
- So we can do:

```
double x, shifter, tmp;
int n;
x = ....
shifter = 6755399441055744.0;           // computation that gives x
tmp = x + shifter;                      // shifter = 2^52 + 2^51
n = (*((unsigned long long int *) &tmp) // tmp = round(2^52 + 2^51 + x)
    & 0xffffffff);                      // n = nearestint(x)
```

Example 2: multiplication by 2^E

- We often need to
 - decompose a floating-point number x into exponent E and significand m
 - reconstruct it knowing E and m
 - or multiply by 2^E
- The construction of 2^E on a variable in double with a variable int E can be done as follows;

```
int E;
double x, twoE;
unsigned long long int tmp;
tmp = E + 1023;           // add double prec. bias to signed exponent
tmp <<= 52;               // shift last exponent bit to the right
twoE = *((double *) &tmp); // make double variable out of it
x = x * twoE;            // e.g. multiply x by  $2^E$ 
```

Our toy exponential

```
// A very crude "toy" implementation of exp(x)
//
// About 45 bits of accuracy. No checks for NaN, Inf whatsoever.
//
double Exp(double x) {
    double z, n, t, r, P, tbl, y;
    uint32_t E, idx, N;
    doubleCaster shiftedN, twoE;

    // Argument reduction
    z = x * TWO_4_RCP_LN_2;           // z = x * 2^4 * 1/ln(2)
    shiftedN.d = z + TWO_52_P_51;     // shiftedN.d = nearestint(z) + 2^52 + 2^51
    n = shiftedN.d - TWO_52_P_51;     // n = nearestint(z) as double
    N = shiftedN.i[LO];              // N = nearestint(z) as integer
    E = N >> 4;                     // E = floor(2^4 * N)
    idx = N & 0x0F;                  // idx = N - E * 2^4
    t = n * TWO_M_4_LN_2;            // t = n * 2^-4 * ln(2)
    r = x - t;                      // r = x - t

    // Polynomial approximation          p(r) approximates exp(r)
    P = c0 + r * (c1 + r * (c2 + r * (c3 + r * (c4 + r * c5))));

    // Table access
    tbl = table[idx];                // tbl = 2^(2^4 * idx)

    // Reconstruction
    twoE.i[HI] = (E + 1023) << 20;
    twoE.i[LO] = 0;                  // twoE.d = 2^E
    y = twoE.d * (tbl * P);         // y = 2^E * tbl * P

    return y;
}
```

Outline

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

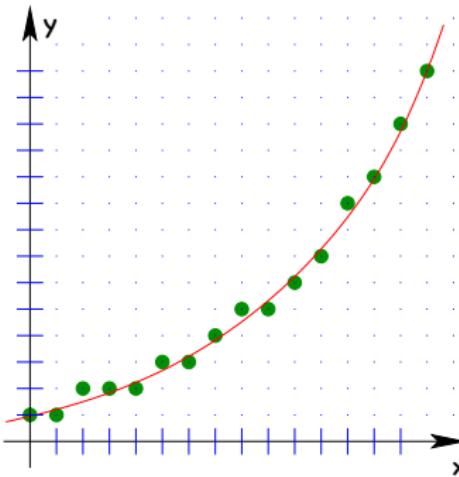
Mathematical functions on a computer

- Developing mathematical functions on computers
 - Programs computing the value $f(x)$ for a given x



- Different algorithms can be used to compute f
- The result of the computation must be the same everywhere.
 - deterministic results
- ⇒ Correctly rounded mathematical functions

Correctly rounded result of a function f



$$\forall x \in \mathbb{Fk}, \quad x \xrightarrow{f} f(x) \xrightarrow{o_k} o_k(f(x))$$

Definition

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a function and $o_k : \mathbb{R} \rightarrow \mathbb{Fk}$ be a rounding.
The function $F : \mathbb{Fk}^n \rightarrow \mathbb{Fk}$ is a correctly rounding of f iff

$$\forall x \in \mathbb{Fk}^n, \quad F(x) = o_k(f(x))$$

Hard cases to round

Correctly rounding $\circ_k(f(x))$:

- Combination of two phenomena
 - The value of $f(x)$, transcendental, cannot be well approximated.
 - The rounding \circ_k changes at the rounding boundaries.



- The value $y = f(x)$ cannot be computed exactly.
- An approximation $\hat{y} = f(x) \cdot (1 + \varepsilon)$ can be computed.
- The precision $\bar{\varepsilon}$ necessary for the worst case is unknown.
- This is the Table Maker's Dilemma.

Accuracy in the worst cases

- Cannot directly compute correct rounding $\circ(f(x))$
- Computation of an approximation $f(x) \cdot (1 + \varepsilon)$ more precise than the relative distance to the worst case
- Use of the following lemma:

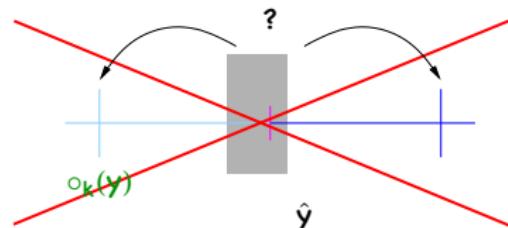
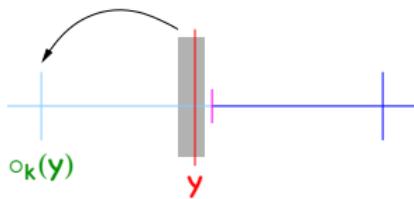
Lemma 1

Let $f = \exp$. and $\circ_{53} : \mathbb{R} \rightarrow \mathbb{F}53$ a rounding in double precision.

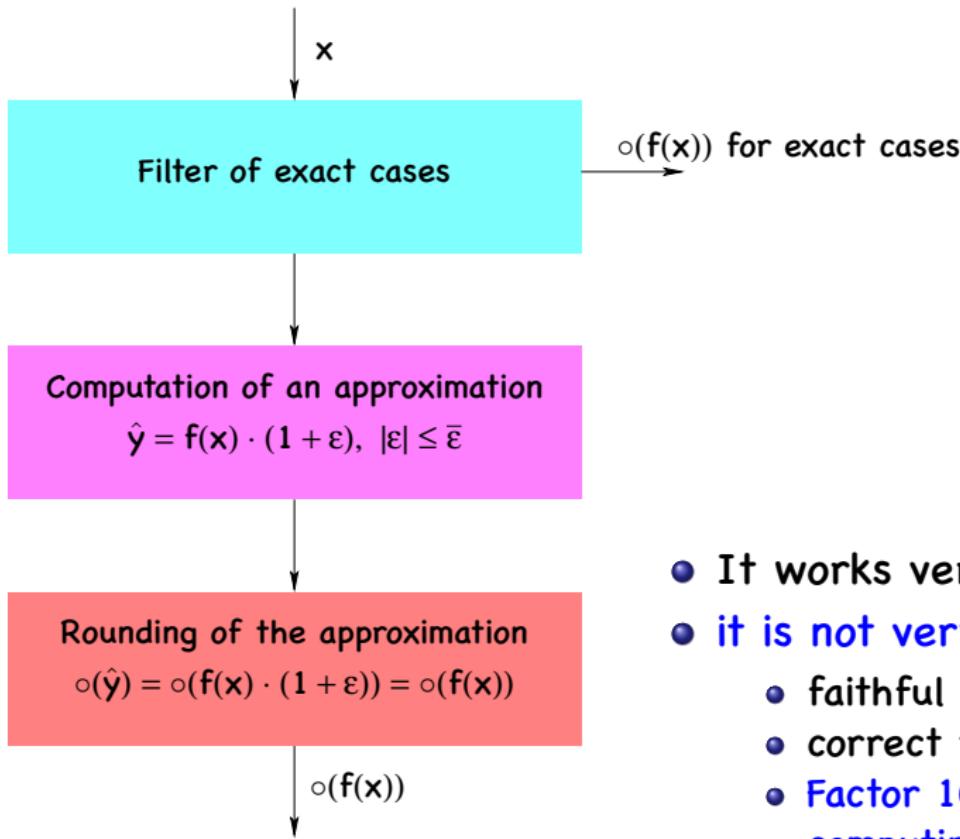
Let $D = \mathbb{F}53 \cap [-744.5; 709]$ be the domain of definition (in double precision) of f .

Then for all $x \in D \setminus \{0\}$ and all $|\varepsilon| \leq 2^{-159}$,

$$\circ_{53}(f(x) \cdot (1 + \varepsilon)) = \circ_{53}(f(x)).$$

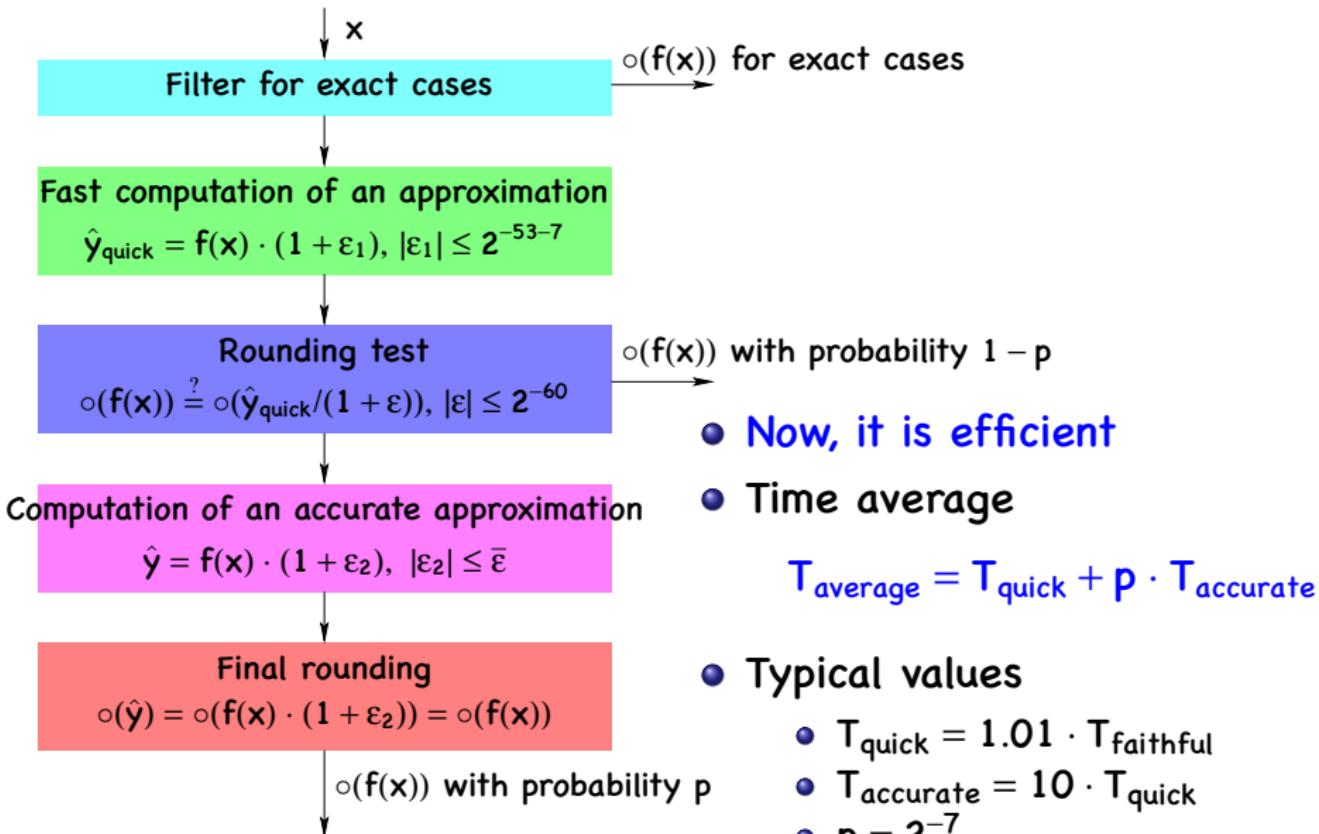


A simple scheme for correct rounding



- It works very well, but...
- it is not very efficient
 - faithful rounding: 53 + 5 bits
 - correct rounding: 159 bits
 - Factor 10 in terms of computing time

Strategy in two steps



Outline

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

Cost of a libm

- Cost of a function in terms of development
 - Beginning of CRLibm: 3-18 person months by function
 - Mean cost CRLibm: 1 person months
 - In industry: 3 person weeks
- Number of functions in a libm
 - IEEE 754-2008: 17 functions
 - Standard libm: around 70 functions
- Maintenance at a high cost
 - Several codes suitable for different architectures
 - Optimization for new architectures
 - Some codes are proved: synchronize proof and code
- Problems with correctly rounded libm
 - Computation of worst cases necessary for each function
 - Codes precisely analyzed and proved:
 - A work for specialists

The problem with proof

- Correct rounding implies a large precision.
- The cases where the maximal precision is needed are very rare.

example : $\exp, \geq 2^{58}$ valid inputs

bits needed	number of cases
159	1
150	6
140	36
130	121
:	:

⇒ a detection of bug by sampling is not sufficient

Toward an automatic generation of libm

- Assembly language is tough \Rightarrow compilers
- A libm is difficult to write \Rightarrow automatic generation!
- MetaLibm¹ – a prototype generator:
 - Input:
 - A function $f : \mathbb{R} \rightarrow \mathbb{R}$
 - A small domain $\text{dom} = [a; b]$, after argument reduction
 - A targeted accuracy $\varepsilon \in \mathbb{R}^+$
 - Output:
 - A code F such that
- $$\forall x \in \text{dom}, \quad \left| \frac{F(x) - f(x)}{f(x)} \right| \leq \varepsilon$$
 - A formal proof for that property
- Ingredients:
 - Polynomial approximation with floating-point coefficients
 - Evaluation in multi-precision arithmetic (multi-double)

Generation of argument reduction

- Generate a code for exp in $[-2^{-5}; 2^{-5}]$ is not sufficient
- \Rightarrow we need to generate codes for argument reduction
- Argument reduction – a major challenge:
 - A multitude of heterogeneous techniques
 - All the remarkable identities have to be coded?
 - How to choose the good one?
 - Codes very irregular
 - Mix of floating-point and fixed-point numbers
 - Tricks for manipulating floating-point numbers
 - Proof by hand difficult to understand
- But: many possibilities then
 - Research for compromise tabulation / approximation
 - Codes generated for any precisions

Outline

- 1 Introduction
- 2 Overview of the implementation of a function
- 3 Polynomial approximation
- 4 Argument reduction
- 5 How to obtain good performances?
- 6 Correctly rounded functions
- 7 Automation
- 8 Conclusion

Conclusion

- The libm functions are based on simple concept:
 - An argument reduction to work on a small domain
 - A polynomial approximation
 - computed by Remez algorithm
 - Some lookup tables
 - A reconstruction
- At first, it seems to be simple
 - because it relies on old result (Taylor, Tchebychev and Remez)
 - But: it is needed to known floating-point arithmetic very well in order to provide efficient code to compute
 - exp in around 40 cycles
- Correctly rounding is possible but difficult to obtain
- An active research area is to write generators of libms.

Some references

Muller,

Elementary Functions, Algorithms and Implementation,
Birkhäuser, 2016.



Abramowitz and Stegun,
Handbook of Mathematical Functions,
National Bureau of Standards.

And also

- P.-T. Tang, Table-driven implementation of the exponential function in IEEE floating-point arithmetic, TOMS, 15(2), 1989.
- E. W. Cheney, Introduction to Approximation Theory, New York, 2nd edition, 1982.