

Introduction to Computer Arithmetic

Floating-point arithmetic and error analysis (AFAE)

Sorbonne Université

Contents

1 Exercise 1: Representation of Signed Integers	2
1.1 Question 1: Three ways to represent an 8-bit signed integer	2
1.1.1 Method 1: Sign-Magnitude	2
1.1.2 Method 2: One's Complement	2
1.1.3 Method 3: Two's Complement (Most Common)	2
2 Exercise 2: IEEE-754 Single Precision Representation	2
2.1 IEEE-754 Single Precision Format	2
2.2 Number Representations	3
2.2.1 1. Number: 13	3
2.2.2 2. Number: 0.4375	3
2.2.3 3. Number: -0.4375	3
2.2.4 4. Number: $1 + 2^{-24}$	3
2.2.5 5. Number: $1 + 2^{-24} - 2^{-25}$	4
2.2.6 6. Number: $1 + 2^{-24} + 2^{-25}$	4
2.2.7 7. Number: $1/7$	4
2.2.8 8. Number: 2^{-130}	4
2.3 Question 2: Product of $a = 4097$ and $b = 8449$	5
3 Exercise 3: Problem with Double Rounding	5
3.1 Given	5
3.2 Question 1: Compute $x + y$ exactly	6
3.3 Question 2: Rounding to double precision, then to single precision	6
3.4 Question 3: Direct rounding from exact sum to single precision	6
3.5 Observation	6
4 Exercise 4: Computation of Square Root and Division	7
4.1 Question 1: Newton-Raphson for Square Root	7
4.2 Question 2: Convergence Rate	8
4.3 Question 3: Division using Newton-Raphson	8
5 Summary	9
5.1 Key Takeaways	9
5.2 Important Formulas	10

1 Exercise 1: Representation of Signed Integers

1.1 Question 1: Three ways to represent an 8-bit signed integer

1.1.1 Method 1: Sign-Magnitude

- **Format:** 1 bit for sign (0=positive, 1=negative) + 7 bits for magnitude
- **Range:** $-(2^7 - 1)$ to $+(2^7 - 1) = -127$ to $+127$
- **19:** 0001 0011 (sign=0, magnitude=19)
- **-19:** 1001 0011 (sign=1, magnitude=19)

1.1.2 Method 2: One's Complement

- **Format:** Positive numbers as usual; negative numbers by flipping all bits
- **Range:** $-(2^7 - 1)$ to $+(2^7 - 1) = -127$ to $+127$
- **19:** 0001 0011
- **-19:** 1110 1100 (flip all bits of 19)

1.1.3 Method 3: Two's Complement (Most Common)

- **Format:** Positive numbers as usual; negative numbers by flipping bits and adding 1
- **Range:** -2^7 to $+(2^7 - 1) = -128$ to $+127$
- **19:** 0001 0011
- **-19:** 1110 1101 (flip bits: 1110 1100, then add 1)

Advantages of Two's Complement

- No duplicate zero
- Simpler arithmetic circuits
- Addition/subtraction use the same hardware

2 Exercise 2: IEEE-754 Single Precision Representation

2.1 IEEE-754 Single Precision Format

S 1 bit	Exponent 8 bits	Mantissa 23 bits
------------	--------------------	---------------------

Formula: $(-1)^S \times 1.M \times 2^{E-127}$

2.2 Number Representations

2.2.1 1. Number: 13

Binary conversion: $13 = 1101_2 = 1.101_2 \times 2^3$

- **Sign (S)**: 0 (positive)
- **Exponent (E)**: $3 + 127 = 130 = 10000010_2$
- **Mantissa (M)**: 101 followed by 20 zeros

IEEE-754: 0 10000010 10100000000000000000000000000000

Hex: 0x41500000

2.2.2 2. Number: 0.4375

Binary conversion:

$$\begin{aligned} 0.4375 \times 2 &= 0.875 \rightarrow 0 \\ 0.875 \times 2 &= 1.75 \rightarrow 1 \\ 0.75 \times 2 &= 1.5 \rightarrow 1 \\ 0.5 \times 2 &= 1.0 \rightarrow 1 \end{aligned}$$

$$0.4375 = 0.0111_2 = 1.11_2 \times 2^{-2}$$

- **Sign (S)**: 0
- **Exponent (E)**: $-2 + 127 = 125 = 01111101_2$
- **Mantissa (M)**: 11 followed by 21 zeros

IEEE-754: 0 01111101 11000000000000000000000000000000

Hex: 0x3EE00000

2.2.3 3. Number: -0.4375

Same as 0.4375 but with sign bit = 1

IEEE-754: 1 01111101 11000000000000000000000000000000

Hex: 0xBEE00000

2.2.4 4. Number: $1 + 2^{-24}$

This is exactly representable in single precision!

$$1 + 2^{-24} = 1.0000000000000000000000000000001_2 \times 2^0$$

- **Sign (S)**: 0
- **Exponent (E)**: $0 + 127 = 127 = 01111111_2$
- **Mantissa (M)**: 23 zeros followed by 1 in the last position

IEEE-754: 0 01111111 00000000000000000000000000000001

Hex: 0x3F800001

2.2.5 5. Number: $1 + 2^{-24} - 2^{-25}$

$$= 1 + 2^{-25}(2 - 1) = 1 + 2^{-25}$$

This cannot be exactly represented (mantissa only has 23 bits). The exact binary would need a 1 in position 25.

Rounding to nearest: This rounds to 1.0 (the 2^{-25} is below the precision)

IEEE-754: 0 01111111 00000000000000000000000000000000

Hex: 0x3F800000

2.2.6 6. Number: $1 + 2^{-24} + 2^{-25}$

$$= 1 + 2^{-25}(2 + 1) = 1 + 3 \times 2^{-25} = 1 + 1.5 \times 2^{-24}$$

This is exactly halfway between $1 + 2^{-24}$ and $1 + 2 \times 2^{-24}$.

Rounding to nearest (even): Rounds to $1 + 2 \times 2^{-24} = 1 + 2^{-23}$

IEEE-754: 0 01111111 00000000000000000000000000000010

Hex: 0x3F800002

2.2.7 7. Number: $1/7$

$$1/7 = 0.142857142857\dots \text{ (repeating)}$$

Converting to binary: $1/7 = 0.001001001001\dots_2$ (repeating pattern: 001)

$$= 1.001001001\dots_2 \times 2^{-3}$$

- **Sign (S):** 0
- **Exponent (E):** $-3 + 127 = 124 = 01111100_2$
- **Mantissa (M):** 00100100100100100100100

IEEE-754: 0 01111100 00100100100100100100100

Hex: 0x3E124925

Warning

Note: $1/7$ cannot be exactly represented - it's a repeating binary fraction

2.2.8 8. Number: 2^{-130}

This is a **denormalized (subnormal) number** because the exponent would be $-130 + 127 = -3$, which is less than 0.

For denormalized numbers:

- Exponent bits = 00000000
- Value = $(-1)^S \times 0.M \times 2^{-126}$

$$2^{-130} = 2^{-126} \times 2^{-4} = 0.0001_2 \times 2^{-126}$$

- **Sign (S):** 0
- **Exponent (E):** 00000000 (denormalized)
- **Mantissa (M):** 0001 followed by 19 zeros

IEEE-754: 0 00000000 00010000000000000000000000000000

Hex: 0x00080000

2.3 Question 2: Product of $a = 4097$ and $b = 8449$

Given:

- $a = 4097 = 2^{12} + 1 = 1.000000000001_2 \times 2^{12}$
 - $b = 8449 = 2^{13} + 2^8 + 1 = 1.0000001000001_2 \times 2^{13}$

Step 1: Represent a in single precision

- Sign: 0
 - Exponent: $12 + 127 = 139 = 10001011_2$
 - Mantissa: 0000000001 + 20 zeros

Step 2: Represent b in single precision

- Sign: 0
 - Exponent: $13 + 127 = 140 = 10001100_2$
 - Mantissa: 0000001000001 + 18 zeros

Step 3: Compute exact product

$$\begin{aligned}
 a \times b &= (2^{12} + 1) \times (2^{13} + 2^8 + 1) \\
 &= 2^{25} + 2^{20} + 2^{12} + 2^{13} + 2^8 + 1 \\
 &= 2^{25} + 2^{20} + 2^{13} + 2^{12} + 2^8 + 1 \\
 &= 34,603,009
 \end{aligned}$$

In binary: 10000100000011000100000001₂

Step 4: Normalize

$$= 1.0000100000011000100000001_2 \times 2^{25}$$

Step 5: Round to 23 bits (single precision)

The mantissa has more than 23 bits. Keeping first 23 bits after the implicit 1: 00001000000110001000000
Looking at bit 24: 0, and bit 25: 1. Since we need to round and bit 24 is 0, we round down
(keep as is).

Result c:

- Sign: 0
 - Exponent: $25 + 127 = 152 = 10011000_2$
 - Mantissa: 00001000000110001000000

IEEE-754: 0 10011000 00001000000110001000000

Hex: 0x4C080620

3 Exercise 3: Problem with Double Rounding

3.1 Given

3.2 Question 1: Compute $x + y$ exactly

Since y is multiplied by 2^{-39} , we need to align the exponents.

x has exponent 0 (assuming normalized form with exponent bits decoded). y needs to be shifted right by 39 positions to align with x .

Exact sum:

$$x = 1.011010100000100111100110111111111110010101111110_2$$

$$y = 0.001101010000000111\dots_2$$

Adding these together:

$$x + y = 1.01101010000010011110011011111111111000010111111011\dots_2$$

The exact result has more than 53 bits of precision in the fractional part.

3.3 Question 2: Rounding to double precision, then to single precision

Step A: Round $x + y$ to double precision (53-bit mantissa)

The exact sum needs to be rounded to 53 bits after the binary point.

Looking at the 53 bits of mantissa: 01101010000010011110011011111111111000010111111011...

The 54th bit is used for rounding. Applying round-to-nearest:

Double precision result: 1.01101010000010011110011011111111111000010111111011₂

After rounding, we keep 53 bits: 01101010000010011110011011111111111000010111111011

Step B: Round this double precision result to single precision (24-bit mantissa)

Now we take the double precision result and round to 23 bits (+ 1 implicit bit).

Mantissa in double: 01101010000010011110011011111111111000010111111011

Keep first 23 bits: 01101010000010011110011

The 24th bit is 0, the 25th bit is 1, and there are more non-zero bits after. Since bit 24 is 0, we round down.

Single precision result from double rounding: 1.01101010000010011110011₂

3.4 Question 3: Direct rounding from exact sum to single precision

Taking the exact sum and rounding directly to single precision (23-bit mantissa):

Exact: 1.01101010000010011110011011111111111000010111111011\dots_2

First 23 bits of mantissa: 01101010000010011110011

Bit 24: 0, Bit 25: 1, Following bits: have 1's

For round-to-nearest with bit 24 = 0, we round down.

Single precision result from direct rounding: 1.01101010000010011110011₂

3.5 Observation

Double Rounding Problem

Key insight: Double rounding can produce different results than direct rounding when:

1. The value is exactly halfway between two representable values at the intermediate precision
2. The tie-breaking rule (round to even) at the intermediate precision affects the final result

In this specific case, the bits after position 53 in the exact sum create a situation where:

- Rounding to double precision first may round up
- Then rounding that double to single may go in a different direction
- Than directly rounding the original to single precision

This demonstrates the “double rounding problem”:

Rounding a value twice (to intermediate then final precision) can give a different result than rounding directly to the final precision.

This is why FPU operations should avoid double rounding when possible.

4 Exercise 4: Computation of Square Root and Division

4.1 Question 1: Newton-Raphson for Square Root

Goal: Compute \sqrt{a}

Method: Find the root of $f(x) = x^2 - a$

Newton-Raphson formula:

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)} \quad (1)$$

Derivatives:

- $f(x) = x^2 - a$
- $f'(x) = 2x$

Iteration formula:

$$\begin{aligned} x_{n+1} &= x_n - \frac{x_n^2 - a}{2x_n} \\ &= x_n - \frac{x_n}{2} + \frac{a}{2x_n} \\ &= \frac{x_n}{2} + \frac{a}{2x_n} \end{aligned}$$

Simplified:

$$x_{n+1} = \frac{1}{2} \left(x_n + \frac{a}{x_n} \right) \quad (2)$$

This is the classic formula for computing square roots!

Example: $\sqrt{2}$ with $x_0 = 1.5$

- $x_1 = \frac{1}{2}(1.5 + 2/1.5) = \frac{1}{2}(1.5 + 1.333\dots) = 1.4166\dots$
- $x_2 = \frac{1}{2}(1.4166\dots + 2/1.4166\dots) = 1.4142\dots$
- $\sqrt{2} \approx 1.41421356\dots$

4.2 Question 2: Convergence Rate

Newton-Raphson has **quadratic convergence**: the number of correct bits approximately doubles at each iteration.

Given: x_0 has 4 bits of accuracy

Iteration	Bits of accuracy
0	4
1	≈ 8
2	≈ 16
3	≈ 32
4	≈ 64

For 24 bits (single precision):

- Start: 4 bits
- After iteration 1: ≈ 8 bits
- After iteration 2: ≈ 16 bits
- After iteration 3: ≈ 32 bits ✓

Answer: 3 iterations needed for 24-bit accuracy

For 53 bits (double precision):

- Start: 4 bits
- After iteration 1: ≈ 8 bits
- After iteration 2: ≈ 16 bits
- After iteration 3: ≈ 32 bits
- After iteration 4: ≈ 64 bits ✓

Answer: 4 iterations needed for 53-bit accuracy

4.3 Question 3: Division using Newton-Raphson

Goal: Compute a/b

Method: Instead of dividing, compute $1/b$ and then multiply by a .

Find root of: $f(x) = 1/x - b = 0$, which means $x = 1/b$

Alternatively, use $f(x) = b - 1/x = 0$

Better formulation: $f(x) = 1 - bx = 0$

- This avoids division in the iteration itself!
- $f'(x) = -b$

Newton-Raphson formula:

$$\begin{aligned} x_{n+1} &= x_n - \frac{1 - bx_n}{-b} \\ &= x_n + \frac{1 - bx_n}{b} \end{aligned}$$

Simplified:

$$x_{n+1} = x_n(2 - bx_n) \quad (3)$$

This formula computes $1/b$ using only multiplications and subtractions!

Final step: Once we have $1/b$, compute $a/b = a \times (1/b)$

Algorithm:

1. Choose initial approximation $x_0 \approx 1/b$
2. Iterate: $x_{n+1} = x_n(2 - bx_n)$ until convergence
3. Result: $a/b \approx a \times x_n$

Example: Compute $7/3$

- Find $1/3$ first using $b = 3$
- $x_0 = 0.3$ (initial guess)
- $x_1 = 0.3(2 - 3 \times 0.3) = 0.3(2 - 0.9) = 0.33$
- $x_2 = 0.33(2 - 3 \times 0.33) = 0.33(2 - 0.99) = 0.3333$
- Then: $7/3 = 7 \times 0.3333\dots \approx 2.333\dots$

Advantage

This method avoids hardware division circuits entirely, using only multiplication and subtraction!

5 Summary

5.1 Key Takeaways

1. **Integer Representation:** Two's complement is the standard due to its simplicity and no duplicate zero
2. **IEEE-754 Floating-Point:** Understanding the bit layout (sign, exponent, mantissa) is crucial for:
 - Recognizing representable numbers
 - Understanding rounding behavior
 - Identifying subnormal numbers
3. **Double Rounding Problem:** Demonstrates that rounding precision matters and can affect final results
4. **Newton-Raphson:**
 - Quadratic convergence (bits double each iteration)
 - Can be adapted for both square root and division
 - Division method avoids hardware division circuits

5.2 Important Formulas

Square Root:

$$x_{n+1} = \frac{1}{2} \left(x_n + \frac{a}{x_n} \right) \quad (4)$$

Division (computing $1/b$):

$$x_{n+1} = x_n(2 - bx_n) \quad (5)$$

Tutorial solved with Claude by Giulia Lionetti - Sorbonne Université